

# Solutions to exercises on Scheduling

1. Explain how the following pairs of scheduling criteria can conflict:

- (a) Maximizing CPU utilization while minimizing response time

sol: CPU utilization is increased if the overheads associated with context switching is minimized. The context switching overheads could be lowered by performing context switches infrequently. This could, however, result in increasing the response time for processes.

- (b) Minimizing average turnaround time while minimizing waiting time

sol: Average turnaround time is minimized by executing the shortest tasks first. Such a scheduling policy could, however, starve long-running tasks and thereby increase their waiting time.

- (c) Maximizing I/O device utilization and maximizing CPU utilization

sol: CPU utilization is maximized by running long-running CPU-bound tasks without performing context switches. I/O device utilization is maximized by scheduling I/O-bound jobs as soon as they become ready to run, thereby incurring the overheads of context switches.

2. “Lottery scheduling” is a scheduling algorithm that works by assigning lottery tickets to threads, which are used for allocating CPU time. Whenever the scheduler needs to schedule a thread, a lottery ticket is chosen at random, and the thread holding that ticket gets the CPU. Assume we have a scheduler that implements the lottery scheduling by holding a lottery 50 times each second, with each lottery winner getting 20 milliseconds of CPU time ( $20 \text{ milliseconds} \times 50 = 1 \text{ second}$ ). Describe how this scheduler can ensure that higher-priority threads receive more attention from the CPU than lower-priority threads.

sol: By assigning more lottery tickets to higher-priority processes.

3. A variation of the round-robin scheduling algorithm is the “regressive round-robin” algorithm. This algorithm assigns each thread a time quantum and a priority. The initial value of a time quantum is 50 milliseconds. However, every time a thread has been allocated the CPU and uses its entire time quantum (does not block for I/O), 10 milliseconds is added to its time quantum, and its priority level is boosted. (The time quantum for a process can be increased to a maximum of 100 milliseconds.) When a process blocks before using its entire time quantum, its time quantum is reduced by 5 milliseconds, but its priority remains the same. What type of thread (CPU-bound or I/O-bound) does the regressive round-robin scheduling algorithm favor? Explain.


sol: favor CPU-bound threads as they are rewarded with a longer time quantum as well as priority boost whenever they consume an entire time quantum. This scheduler does not penalize I/O-bound processes as they are likely to block for I/O before consuming their entire time quantum, but their priority remains the same.

4. Consider the following set of processes, with the length of the CPU burst time given in milliseconds

| <u>Process</u> | <u>Burst Time</u> | <u>Priority</u> |
|----------------|-------------------|-----------------|
| $P_1$          | 2                 | 2               |
| $P_2$          | 1                 | 1               |
| $P_3$          | 8                 | 4               |
| $P_4$          | 4                 | 2               |
| $P_5$          | 5                 | 3               |

The processes are assumed to have arrived in the order  $P_1, P_2, P_3, P_4, P_5$  all at time 0.

- (a) Draw four Gantt charts that illustrate the execution of these processes using the following scheduling algorithms: FCFS, SJF, nonpreemptive priority (a smaller priority number implies a higher priority), and RR (quantum = 1).  
sol:



- (b) What is the turnaround time of each process for each of the scheduling algorithms in part a?

|           | FCFS | Turnaround time |     | Priority |
|-----------|------|-----------------|-----|----------|
|           |      | RR              | SJF |          |
| <i>P1</i> | 10   | 19              | 19  | 16       |
| <i>P2</i> | 11   | 2               | 1   | 1        |
| <i>P3</i> | 13   | 7               | 4   | 18       |
| <i>P4</i> | 14   | 4               | 2   | 19       |
| <i>P5</i> | 19   | 14              | 9   | 6        |

- (c) What is the waiting time of each process for each of these scheduling algorithms?

|           | Waiting time (turnaround time minus burst time) |    |     |          |
|-----------|---|----|-----|----------|
|           | FCFS  | RR | SJF | Priority |
| <i>P1</i> | 0   | 9  | 9   | 6        |
| <i>P2</i> | 10  | 1  | 0   | 0        |
| <i>P3</i> | 11  | 5  | 2   | 16       |
| <i>P4</i> | 13  | 3  | 1   | 18       |
| <i>P5</i> | 14  | 9  | 4   | 1        |

- (d) Which of the algorithms results in the minimum average waiting time (over all processes)?

sol: Shortest Job First

5. The following processes are being scheduled using a preemptive, round-robin scheduling algorithm. Each process is assigned a numerical priority, with a higher number

indicating a higher relative priority. In addition to the processes listed below, the system also has an idle task (which consumes no CPU resources and is identified as  $P_{idle}$ ). This task has priority 0 and is scheduled whenever the system has no other available processes to run. The length of a time quantum is 10 units. If a process is preempted by a higher-priority process, the preempted process is placed at the end of the queue.

| Thread | Priority | Burst | Arrival |
|--------|----------|-------|---------|
| $P_1$  | 40       | 20    | 0       |
| $P_2$  | 30       | 25    | 25      |
| $P_3$  | 30       | 25    | 30      |
| $P_4$  | 35       | 15    | 60      |
| $P_5$  | 5        | 10    | 100     |
| $P_6$  | 10       | 10    | 105     |

- (a) Show the scheduling order of the processes using a Gantt chart.
  - (b) What is the turnaround time for each process?  
 sol: p1: 20-0 = 20, p2: 80-25 = 55, p3: 90 - 30 = 60, p4: 75-60 = 15, p5: 120-100 = 20, p6: 115-105 = 10
  - (c) What is the waiting time for each process?  
 sol: p1: 0, p2: 40, p3: 35, p4: 0, p5: 10, p6: 0
  - (d) What is the CPU utilization rate?  
 sol:  $105/120 = 87.5$  percent.
6. Which of the following scheduling algorithms could result in starvation? a) First-come, first-served; b) Shortest job first; c) Round robin; d) Priority  
 sol: Shortest job first and priority-based scheduling algorithms could result in starvation.
  7. Consider a preemptive priority scheduling algorithm based on dynamically changing priorities. Larger priority numbers imply higher priority. When a process is waiting for the CPU (in the ready queue, but not running), its priority changes at a rate  $\alpha$ ; when it is running, its priority changes at a rate  $\beta$ . All processes are given a priority of 0 when they enter the ready queue. The parameters  $\alpha$  and  $\beta$  can be set to give many different scheduling algorithms.

- (a) What is the algorithm that results from  $\beta > \alpha > 0$ ? sol: FCFS
  - (b) What is the algorithm that results from  $\alpha < \beta < 0$ ? LIFO
8. Explain the differences in how much the following scheduling algorithms discriminate in favor or against short processes (sum of CPU and I/O bursts is short):
- (a) FCFS: sol: FCFS discriminates against short jobs since any short jobs arriving after long jobs will have a longer waiting time.
  - (b) RR: sol: RR treats all jobs equally (giving them equal bursts of CPU time) so short jobs will be able to leave the system faster since they will finish first.
  - (c) Multilevel feedback queues: sol: Multilevel feedback queues work similar to the RR algorithm they discriminate favorably toward short jobs.
9. Consider the scheduling algorithm in the Solaris operating system for time-sharing threads:
- (a) What is the time quantum (in milliseconds) for a thread with priority 10? With priority 55? sol: 160 and 40
  - (b) Assume a thread with priority 35 has used its entire time quantum without blocking. What new priority will the scheduler assign this thread? sol: 35
  - (c) Assume a thread with priority 35 blocks for I/O before its time quantum has expired. What new priority will the scheduler assign this thread? sol: 54
10. Using the Windows scheduling algorithm, determine the numeric priority of each of the following threads
- (a) A thread in the `REALTIME_PRIORITY_CLASS` with a relative priority of `HIGHEST`. sol 26
  - (b) A thread in the `NORMAL_PRIORITY_CLASS` with a relative priority of `NORMAL`. sol 8
  - (c) A thread in the `HIGH_PRIORITY_CLASS` with a relative priority of `ABOVE_NORMAL`. sol 14