

Some exercises on Scheduling

1. Explain how the following pairs of scheduling criteria can conflict:
 - (a) Maximizing CPU utilization while minimizing response time
 - (b) Minimizing average turnaround time while minimizing waiting time
 - (c) Maximizing I/O device utilization and maximizing CPU utilization
2. “Lottery scheduling” is a scheduling algorithm that works by assigning lottery tickets to threads, which are used for allocating CPU time. Whenever the scheduler needs to schedule a thread, a lottery ticket is chosen at random, and the thread holding that ticket gets the CPU. Assume we have a scheduler that implements the lottery scheduling by holding a lottery 50 times each second, with each lottery winner getting 20 milliseconds of CPU time ($20 \text{ milliseconds} \times 50 = 1 \text{ second}$). Describe how this scheduler can ensure that higher-priority threads receive more attention from the CPU than lower-priority threads.
3. A variation of the round-robin scheduling algorithm is the “regressive round-robin” algorithm. This algorithm assigns each thread a time quantum and a priority. The initial value of a time quantum is 50 milliseconds. However, every time a thread has been allocated the CPU and uses its entire time quantum (does not block for I/O), 10 milliseconds is added to its time quantum, and its priority level is boosted. (The time quantum for a process can be increased to a maximum of 100 milliseconds.) When a process blocks before using its entire time quantum, its time quantum is reduced by 5 milliseconds, but its priority remains the same. What type of thread (CPU-bound or I/O-bound) does the regressive round-robin scheduling algorithm favor? Explain.
4. Consider the following set of processes, with the length of the CPU burst time given in milliseconds

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	2	2
P_2	1	1
P_3	8	4
P_4	4	2
P_5	5	3

The processes are assumed to have arrived in the order P_1, P_2, P_3, P_4, P_5 all at time 0.

- (a) Draw four Gantt charts that illustrate the execution of these processes using the following scheduling algorithms: FCFS, SJF, nonpreemptive priority (a smaller priority number implies a higher priority), and RR (quantum = 1).
 - (b) What is the turnaround time of each process for each of the scheduling algorithms in part a?
 - (c) What is the waiting time of each process for each of these scheduling algorithms?
 - (d) Which of the algorithms results in the minimum average waiting time (over all processes)?
5. The following processes are being scheduled using a preemptive, round-robin scheduling algorithm. Each process is assigned a numerical priority, with a higher number indicating a higher relative priority. In addition to the processes listed below, the system also has an idle task (which consumes no CPU resources and is identified as P_{idle}). This task has priority 0 and is scheduled whenever the system has no other available processes to run. The length of a time quantum is 10 units. If a process is preempted by a higher-priority process, the preempted process is placed at the end of the queue.

Thread	Priority	Burst	Arrival
P_1	40	20	0
P_2	30	25	25
P_3	30	25	30
P_4	35	15	60
P_5	5	10	100
P_6	10	10	105

- (a) Show the scheduling order of the processes using a Gantt chart.
 - (b) What is the turnaround time for each process?
 - (c) What is the waiting time for each process?
 - (d) What is the CPU utilization rate?
6. Which of the following scheduling algorithms could result in starvation? a) First-come, first-served; b) Shortest job first; c) Round robin; d) Priority
7. Consider a preemptive priority scheduling algorithm based on dynamically changing priorities. Larger priority numbers imply higher priority. When a process is waiting for the CPU (in the ready queue, but not running), its priority changes at a rate α ; when it is running, its priority changes at a rate β . All processes are given a priority of 0 when they enter the ready queue. The parameters α and β can be set to give many different scheduling algorithms.
- (a) What is the algorithm that results from $\beta > \alpha > 0$?

- (b) What is the algorithm that results from $\alpha < \beta < 0$?
- 8. Explain the differences in how much the following scheduling algorithms discriminate in favor or against short processes (sum of CPU and I/O bursts is short):
 - (a) FCFS:
 - (b) RR:
 - (c) Multilevel feedback queues:
- 9. Consider the scheduling algorithm in the Solaris operating system for time-sharing threads:
 - (a) What is the time quantum (in milliseconds) for a thread with priority 10? With priority 55?
 - (b) Assume a thread with priority 35 has used its entire time quantum without blocking. What new priority will the scheduler assign this thread?
 - (c) Assume a thread with priority 35 blocks for I/O before its time quantum has expired. What new priority will the scheduler assign this thread?
- 10. Using the Windows scheduling algorithm, determine the numeric priority of each of the following threads
 - (a) A thread in the `REALTIME_PRIORITY_CLASS` with a relative priority of `HIGHEST`.
 - (b) A thread in the `NORMAL_PRIORITY_CLASS` with a relative priority of `NORMAL`.
 - (c) A thread in the `HIGH_PRIORITY_CLASS` with a relative priority of `ABOVE_NORMAL`.