

Artificial Intelligence (IT3160E)

Than Quang Khoat

khoattq@soict.hust.edu.vn

School of Information and Communication Technology
Hanoi University of Science and Technology

2022

Content:

- Introduction of Artificial Intelligence
- Intelligent agent
- **Problem solving: Search**, Constraint satisfaction
 - **Informed search**
- Logic and reasoning
- Knowledge representation
- Machine learning

Reminder: Tree-based search

```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem] applied to STATE(node) succeeds return node
    fringe ← INSERTALL(EXPAND(node, problem), fringe)
```

- A search strategy (i.e., method)
= A way of determining the order to examine the tree's nodes

Informed search

- Uninformed search strategies use only the information contained in the problem definition
 - Not suitable for many practical problems (due to high cost of time and memory)
- Informed search strategies use the *problem-specific knowledge* → The search process is more efficient
 - Best-first search algorithms (Greedy best-first, A*)
 - Local search algorithms (Hill-climbing, Simulated annealing, Local beam)

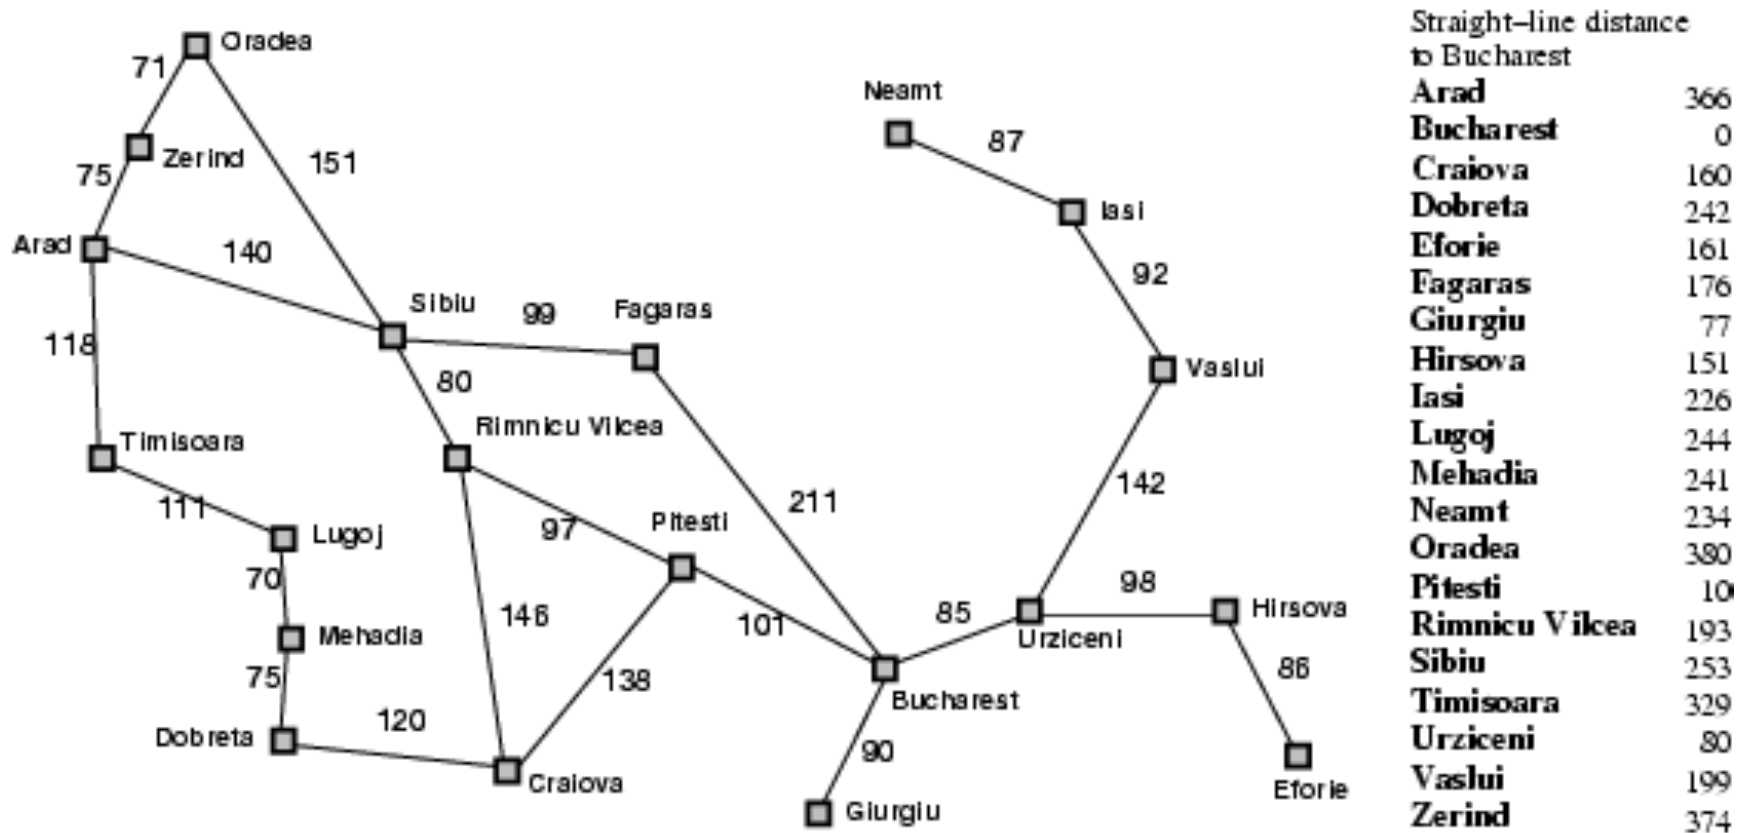
Best-first search

- **Intuitive idea:** use an *evaluation function* $f(n)$ for every node of the search tree
 - To evaluate the "suitability" of that node
 - During the search process, the nodes with the highest suitability are given priority
- **Implementation**
 - Order the nodes in *fringe* in descending order of suitability
- **Best-first search algorithms**
 - Greedy best-first search
 - A* search

Greedy best-first search

- The evaluation function $f(n)$ is a *heuristic* function $h(n)$
- The heuristic function $h(n)$ estimates the **cost from n to goal**
- Example: In the problem of finding a way from Arad to Bucharest: $h_{SLD}(n)$ = Estimated straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal

Greedy best-first search – Example (1)



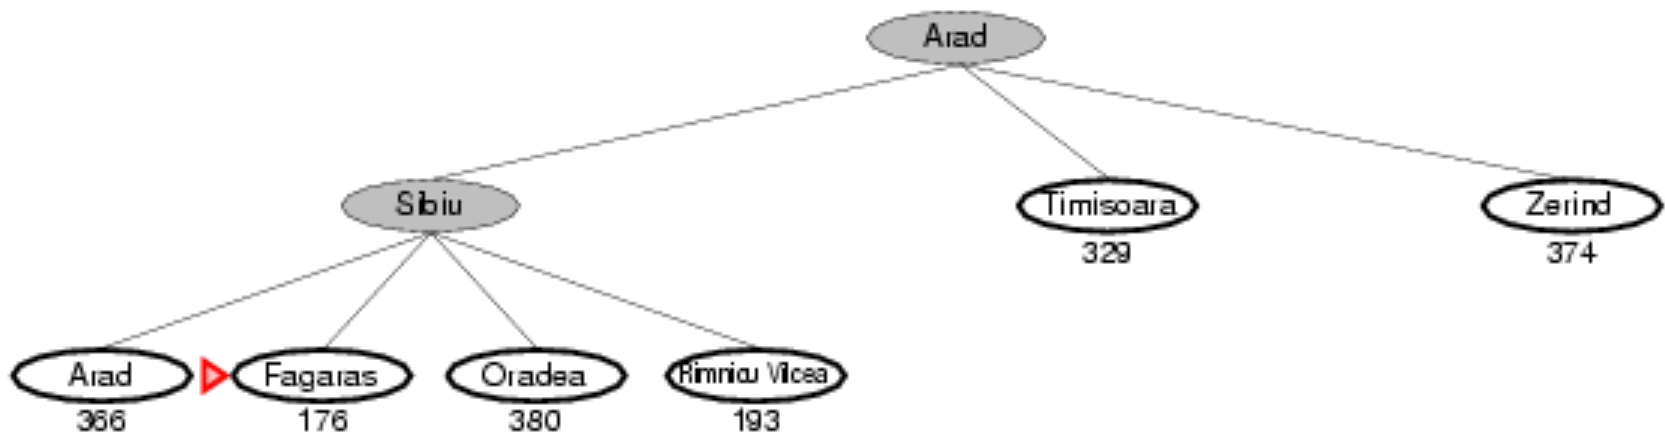
Greedy best-first search – Example (2)



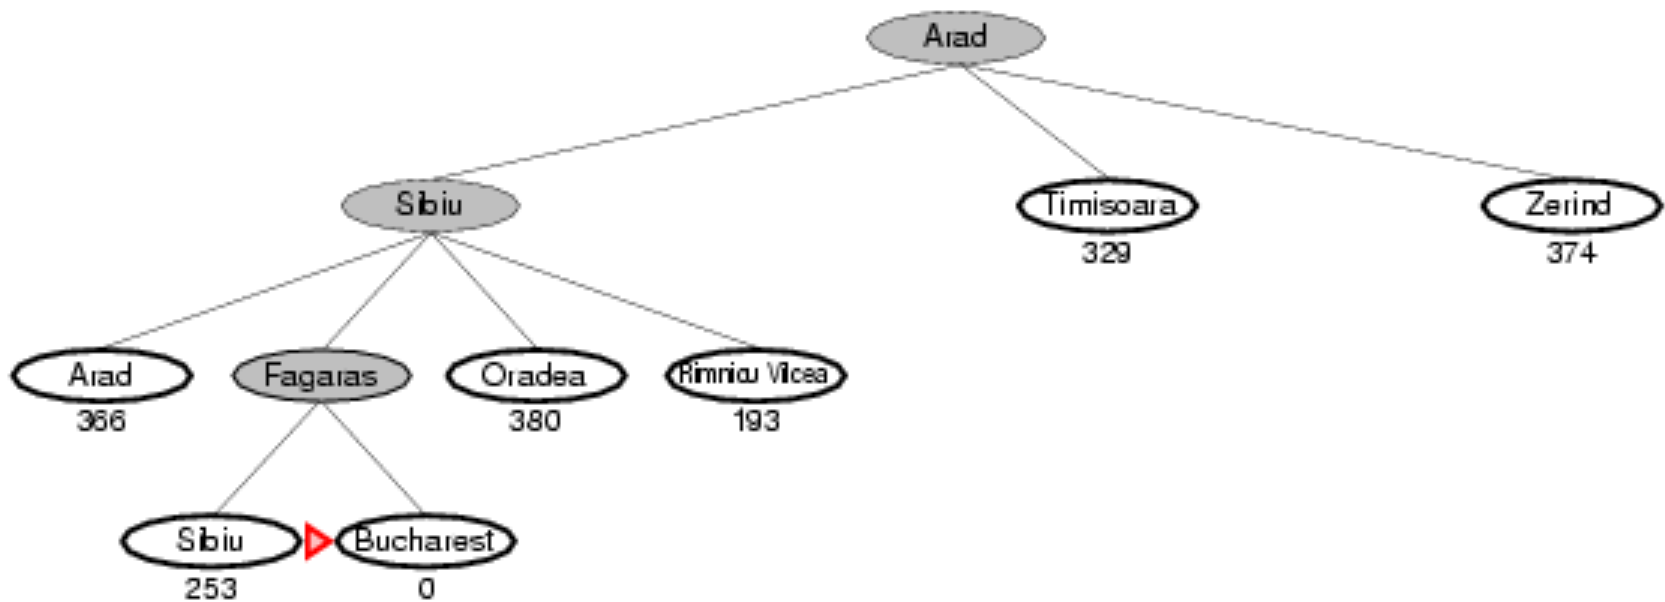
Greedy best-first search – Example (3)



Greedy best-first search – Example (4)



Greedy best-first search – Example (5)



Properties of Greedy best-first search

■ Complete?

- No, because it can get stuck in loops, e.g., lasi → Neamt → lasi → Neamt → ...

■ Time?

- $O(b^m)$
- But a good heuristic function can give dramatic improvement

■ Space?

- $O(b^m)$ – Keeps all nodes in memory


■ Optimal?

- No

A* search

- **Intuitive idea:** Avoid expanding paths that are already (i.e., up to the current moment) determined expensive
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = The cost from the root node to (the current one) n
 - $h(n)$ = The estimated cost from n to goal
 - $f(n)$ = The estimated total cost of path through n to goal


A^* search: Example (1)



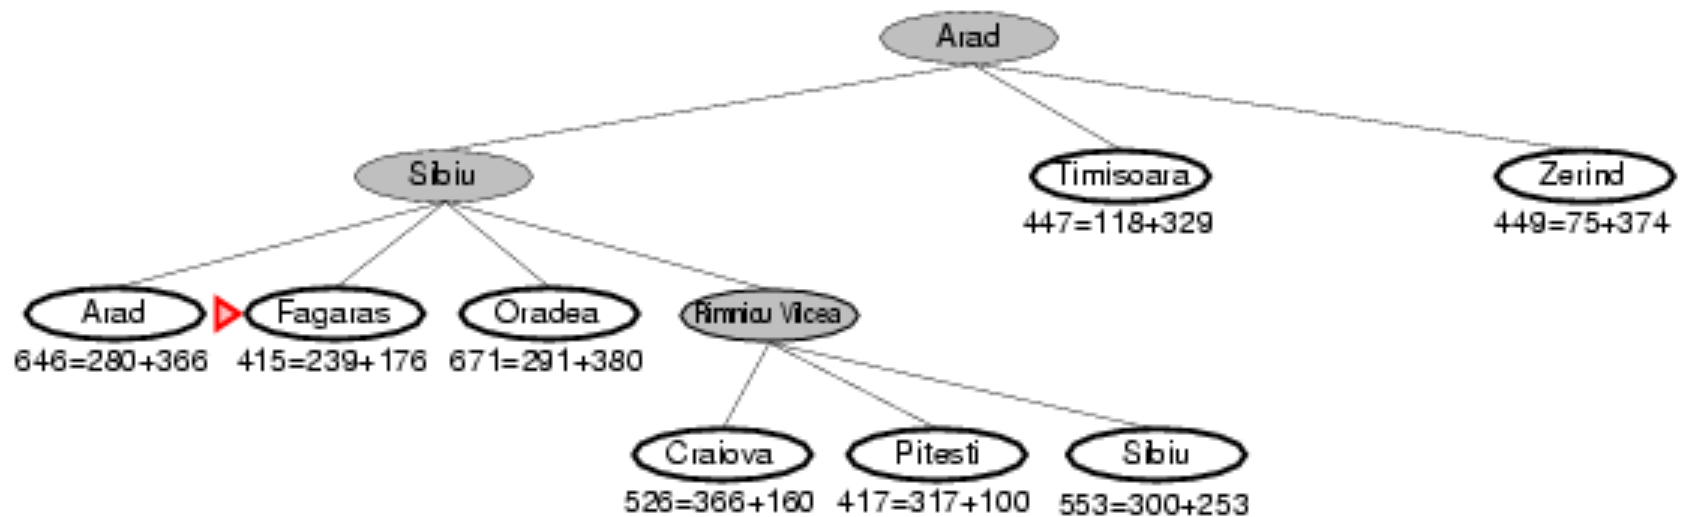
A^* search: Example (2)



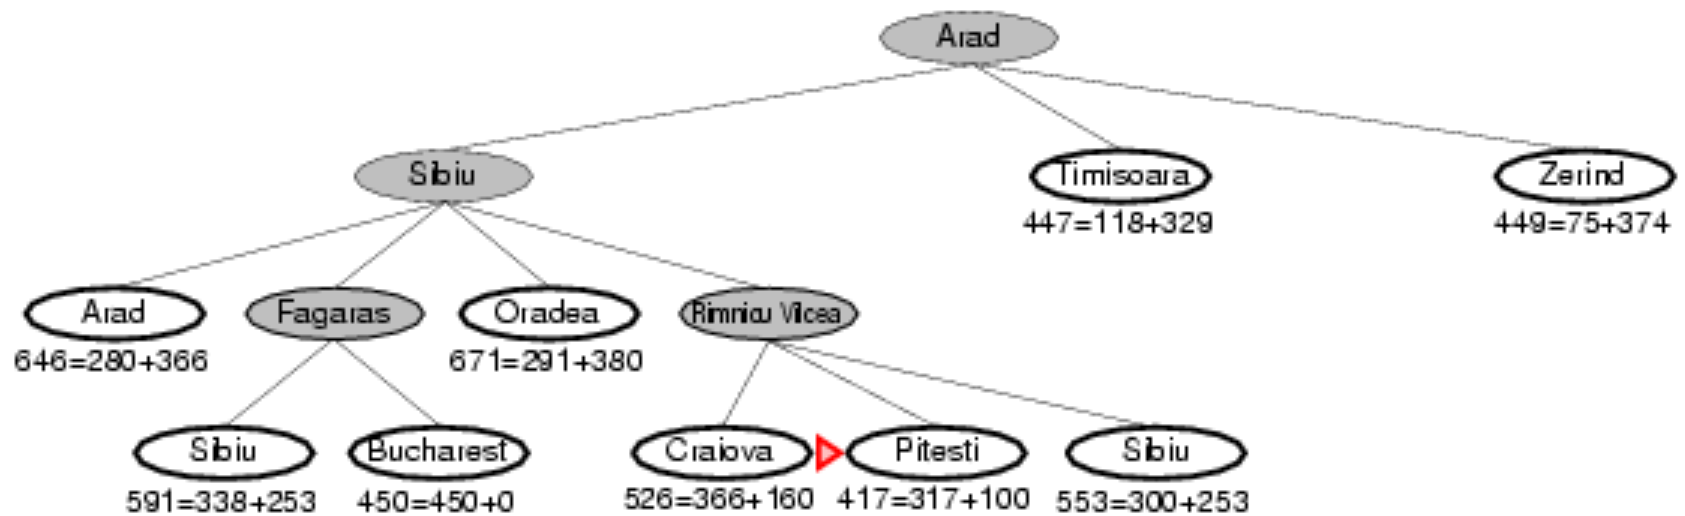
A* search: Example (3)



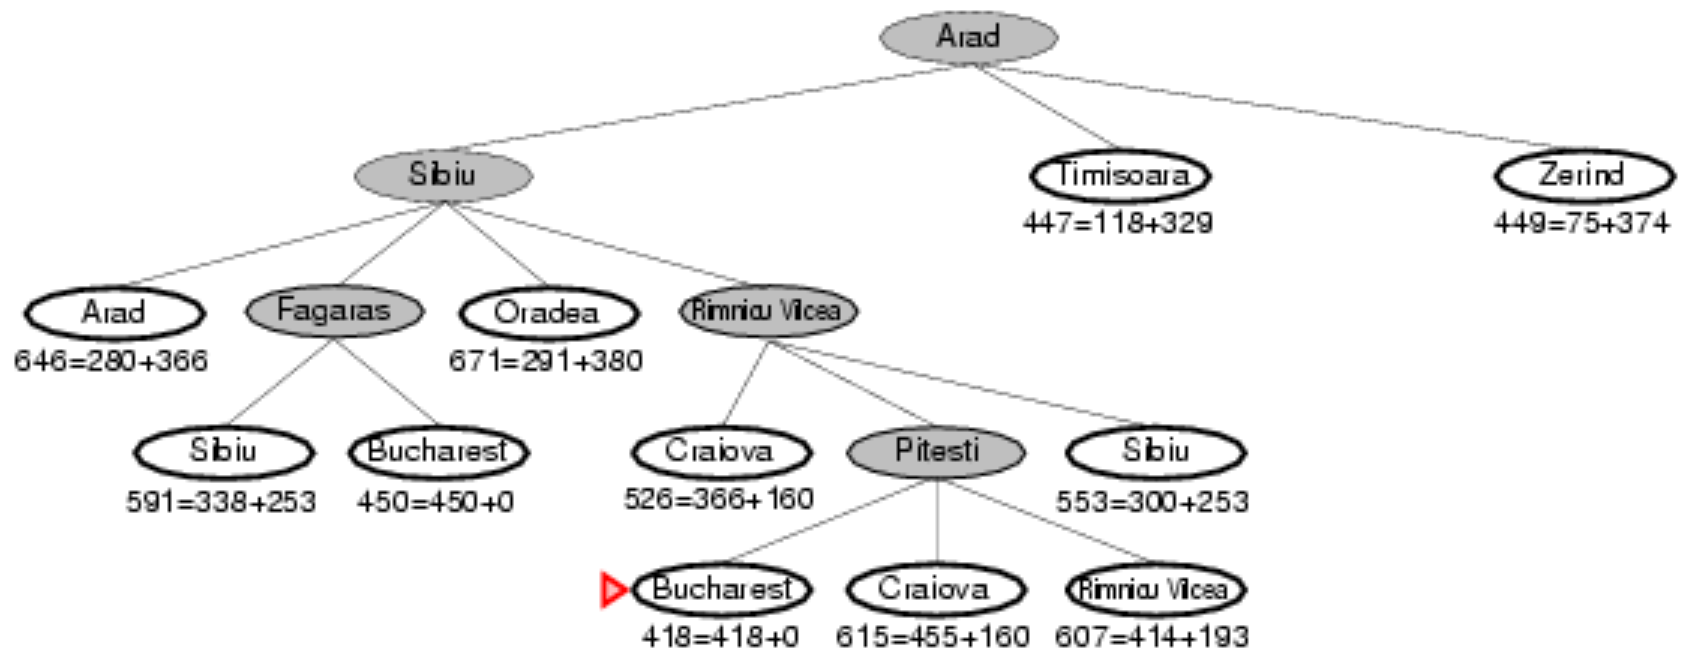
A* search: Example (4)



A* search: Example (5)



A* search: Example (6)




A* search: Properties


- If *the state space is finite* and there is *a solution to avoid repeating the states*, then the A* algorithm is complete (i.e., can find the solution), but the optimal is not guaranteed
- If *the state space is finite* and there is *no solution to avoid repeating the states*, then the A* algorithm is incomplete (i.e., no guarantee to find a solution)
- If *the state space is infinite*, then the A* algorithm is incomplete (i.e., no guarantee to find a solution)
- When is A* optimal?

A* vs. UCS

- Uniform-cost search (UCS) expands in all directions



- A* expands mainly towards the goal, but the optimal is guaranteed




Admissible heuristics

- A heuristic $h(n)$ is admissible if $0 \leq h(n) \leq h^*(n)$ for every node n , where $h^*(n)$ is the true cost to reach to the goal state from n
- An admissible heuristic *never overestimates* the cost to reach the goal
 - It is *optimistic*
- Example: The heuristic $h_{SLD}(n)$ *underestimates* the actual road distance
- **Theorem:** If $h(n)$ is admissible, A^* using TREE-SEARCH is optimal


Optimality of A*: Proof (1)

- Suppose some suboptimal goal G_2 has been generated and is in the *fringe*. Let n be an unexpanded node in the *fringe* such that n is on a shortest path to an optimal goal G



- We have: 1) $f(G_2) = g(G_2)$ since $h(G_2) = 0$
- We have: 2) $g(G_2) > g(G)$ since G_2 is suboptimal
- We have: 3) $f(G) = g(G)$ since $h(G) = 0$
- From 1)+2)+3) we have: 4) $f(G_2) > f(G)$

Optimality of A*: Proof (2)



- We have: 5) $h(n) \leq h^*(n)$ since h is admissible
- From 5) we have: 6) $g(n) + h(n) \leq g(n) + h^*(n)$
- We have: 7) $g(n) + h^*(n) = f(G)$ since n is in the path to G
- From 6)+7) we have: 8) $f(n) \leq f(G)$
- From 4)+8) we have: $f(G_2) > f(n)$ A* never selects G_2 for expansion

Admissible heuristics (1)

Example: For the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = the least number of moves ($\leftarrow, \rightarrow, \uparrow, \downarrow$) to move the misplaced tiles to their correct position

- $h_1(S) = ?$

- $h_2(S) = ?$

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Admissible heuristics (2)

Example: For the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = the least number of moves ($\leftarrow, \rightarrow, \uparrow, \downarrow$) to move the misplaced tiles to their correct position

- $h_1(S) = 8$

- $h_2(S) = 3+1+2+2+2+3+3+2 = 18$

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Dominant heuristic

- Heuristic h_2 **dominates** heuristic h_1 if:
 - Both $h_1(n)$ and $h_2(n)$ are admissible, and
 - $h_2(n) \geq h_1(n)$ for every node n
- If heuristic h_2 dominates heuristic h_1 , then h_2 is better (to be used) for search
- The 8-puzzle: *Search cost = Average number of nodes expanded*
 - For the depth $d = 12$
 - IDS (Iterative Deepening Search): 3,644,035 nodes
 - A*(using heuristic h_1): 227 nodes
 - A*(using heuristic h_2): 73 nodes
 - For the depth $d = 24$
 - IDS (Iterative Deepening Search): Too many nodes
 - A*(using heuristic h_1): 39,135 nodes
 - A*(using heuristic h_2): 1,641 nodes


Consistent heuristics

- A heuristic h is **consistent** if for every node n and every successor n' of n (generated by action a):

$$h(n) \leq c(n,a,n') + h(n')$$

- If heuristic h is consistent, we have:

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n,a,n') + h(n') \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$



That means: $f(n)$ is non-decreasing along any path passing through n

- **Theorem:** If $h(n)$ is consistent, then A^* using GRAPH-SEARCH is optimal

Properties of A^*

■ Complete?

- Yes (Unless there are infinitely many nodes with $f \leq f(G)$)

■ Time?

- Exponential (The number of considered nodes is an exponential function of the solution's path length)

■ Space?

- Keeps all nodes in memory

■ Optimal?

- Yes, for some special conditions

Local search algorithms

- In many optimization problems, the path to the goal is irrelevant
 - The goal state itself = The solution
- The state space = A set of "complete" **configurations**
- **Goal:** *Find a configuration satisfying all the constraints*
 - Example: The n -queens problem (i.e., arrange n queens on a board of size $n \times n$ so that they do not attack each other)
- In such problems, we can use local search algorithms
- Save only a single "current" state (i.e., configuration) at a time
 - Idea:* Try to "improve" this current state using a (predefined) criterion

Example: n -queens problem

- Arrange n ($=4$) queens on a board of size $n \times n$ with no two queens on the same row, column or diagonal



Hill-climbing search: Algorithm

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                     neighbor, a node


  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```


Hill-climbing search: n -queens problem (1)

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18

- Heuristic h = The number of pairs of queens attacking each other, either directly or indirectly
- For the above state: $h = 17$


Hill-climbing search: n -queens problem (2)



- The above state is a local minimum solution
 - With $h = 1$ (i.e., there is still a pair of queens attacking each other)

Hill-climbing search: **issue**

- Problem: Depending on initial state, can get stuck in local optimal
 - Cannot return the global optimal solution



Simulated annealing search

- *Annealing process*: The metal cools and freezes into a crystallized structure
- Simulated annealing search may avoid returning local optima
- Simulated annealing search uses **random search strategy**, which accepts changes that increase the value of the target function (need to maximize) and *also accepts (but limited) changes that decrease* the value of the target function
- Simulated annealing search uses a control parameter T (as in temperature systems)
 - T is a high value at the beginning of the search, and then decreases gradually to 0

Simulated annealing search: Algorithm

- Intuitive idea: Escape local maxima by allowing some "bad" moves but *gradually decrease* their frequency

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to "temperature"
  local variables: current, a node
                   next, a node
                   T, a "temperature" controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to  $\infty$  do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```

Simulated annealing search: Properties

- If T (i.e., that defines the degree of frequency reduction for “bad moves”) decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1

Local beam search

- At a time in the search process, keep track of k – rather than just 1 – best states
- At the beginning of the search: Select k states randomly
- At each iteration, generate all the successors of all k states
- If any one is a goal state, stop (successfully); else select the k best successors from the complete list and repeat

Adversarial search

- IDS and A* only consider the search problems with one agent
- How about an environment with two agents which may have **conflict of interest**?
 - Adversarial search (Tìm kiếm có đối thủ)
- Adversarial search is often used in games

Issues of search in games

- Hard to predict the reaction of the opponent
 - Need to determine a suitable move for each reaction (or move) of the opponent
- Time limit (time-counting games)
 - Difficult (or unable) to find an optimal solution → approximation
- Adversarial search often requires effectiveness (quality of each move and time cost)
→ a hard requirement
- In a zero-sum adversarial game:
 - Winner $>$ Loser
 - Winning score of the winner = Losing score of the loser


Tic Tac Toe (cờ ca-rô)

- This is an adversarial game
 - E.g.: <http://www.ourvirtualmall.com/tictac.htm>
- It consists of two players (e.g., MAX and MIN)
 - Each will move just after the other's move
 - Game termination: Winner will have bonus, while loser will be penalized

Representing an adversarial game

- Components for representing a game
 - *Initial state*: State of the game + Who will move first
 - *Successor function*: return some information (given a move, states)
 - All the admissible moves
 - New state (after the move)
 - *Terminal test*
 - *Utility function* to evaluate each state
- Initial state + admissible moves = Game tree

Game tree for Tic Tac Toe




Optimal strategies

- An *optimal strategy* is a sequence of moves to achieve the goal (e.g., winner)
- The strategy of MAX can depend on the moves by MIN, and vice versa
- MAX needs choosing a strategy that *maximizes* its objective function, assuming that MIN uses optimal moves
 - MIN needs choosing a strategy that *minimizes* its objective function
- This strategy can be determined by considering the MINIMAX value at each node in a game tree

MINIMAX value

- MAX chooses a move with *maximal* MINIMAX value (to maximize its objective function)
- In contrast, MIN chooses a move with *minimal* MINIMAX value (to minimize its objective function)



MINIMAX algorithm

function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*


MINIMAX algorithm: properties

- Completeness
 - Yes (if the game tree is finite)
- Optimality
 - Yes (if the players always choose the optimal move at each step)
- Time complexity
 - $O(b^m)$
- Memory complexity
 - $O(bm)$ (based on DFS)
- For Chess, branching factor $b \approx 35$ and tree depth $m \approx 100$
 - Too expensive – cannot find an optimal strategy


Pruning (cắt tỉa)

- **Issue:** MINIMAX algorithm may have an exponential number of moves to be considered → may not be practical
- We can prune some branches in the tree
- α - β pruning (Alpha-beta pruning):
 - *Idea:* if a branch cannot improve the objective function, we can ignore it!
 - Pruning a bad branch may not affect the solution.


α - β pruning: example (1)




α - β pruning: example (2)




α - β pruning: example (3)



α - β pruning: example (4)




α - β pruning: example (5)



Why calling α - β pruning?

- α is the best move's value of MAX until now at the current branch
- If v is worse than α , MAX will ignore the moves with value v
 - Prune the branches with value v
- β has the same meaning for MIN



α - β pruning algorithm (1)

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for a, s in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return v

α - β pruning algorithm (2)

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to state
            $\beta$ , the value of the best alternative for MIN along the path to state

  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

α - β pruning

- For games with a large state space, α - β pruning is still not good
 - The pruned space is still large
- Domain knowledge about the game can be used to reduce the search space
 - Such a knowledge can enable us to evaluate each state
 - Such an additional knowledge plays a similar role with heuristic function $h(n)$ in the A* algorithm