

Machine Learning

(IT3190E)

Quang Nhat NGUYEN

quang.nguyennhat@hust.edu.vn

Hanoi University of Science and Technology
School of Information and Communication Technology
Academic year 2020-2021

The course's content:

- Introduction
- Performance evaluation of ML system
- **Supervised learning**
 - **Classification problem**
 - **Nearest neighbor learning**
- Unsupervised learning
- Ensemble learning
- Reinforcement learning

Classification problem

- Classification problem belongs to supervised learning
- The goal of the classification problem is to predict a discrete (i.e., nominal) value

$$f: X \rightarrow Y$$

where Y is a finite set of discrete values

Classification problem: Performance evaluation

$$Accuracy = \frac{1}{|D_{test}|} \sum_{x \in D_{test}} Identical(o(x), c(x));$$

$$Identical(a, b) = \begin{cases} 1, & \text{if } (a = b) \\ 0, & \text{if otherwise} \end{cases}$$

- x : A test example in the test set D_{test}
- $o(x)$: The class label produced by the system for the example x
- $c(x)$: The true (desired) class label for the example x

Confusion matrix

- Also called Contingency Table
- **Can be used only for a classification problem**
 - *Cannot be used for a regression problem*
- **TP_i** : The number of examples of class c_i are correctly classified
- **FP_i** : The number of examples not belonging to class c_i are incorrectly classified in class c_i
- **TN_i** : The number of examples not belonging to class c_i are correctly classified
- **FN_i** : The number of examples of class c_i are incorrectly classified into classes different from c_i

(For Class c_i)		Classified by the system	
		Class c_i	Not Class c_i
True class	Class c_i	TP_i	FN_i
	Not Class c_i	FP_i	TN_i

Precision and Recall (1)

- Very often used in evaluation of text mining and information retrieval systems

- **Precision** for class c_i

→ The number of examples correctly classified to class c_i divides the number of examples classified to class c_i

$$\text{Precision}(c_i) = \frac{TP_i}{TP_i + FP_i}$$

- **Recall** for class c_i

→ The number of examples correctly classified to class c_i divides the number of examples of class c_i

$$\text{Recall}(c_i) = \frac{TP_i}{TP_i + FN_i}$$

Precision and Recall (2)

- How to compute the overall Precision and Recall values for all the class labels $C=\{c_i\}$?

- Micro-averaging

$$Precision = \frac{\sum_{i=1}^{|C|} TP_i}{\sum_{i=1}^{|C|} (TP_i + FP_i)}$$

$$Recall = \frac{\sum_{i=1}^{|C|} TP_i}{\sum_{i=1}^{|C|} (TP_i + FN_i)}$$

- Macro-averaging

$$Precision = \frac{\sum_{i=1}^{|C|} Precision(c_i)}{|C|}$$

$$Recall = \frac{\sum_{i=1}^{|C|} Recall(c_i)}{|C|}$$

F_1 measure

- The F_1 evaluation metric is a combination of *Precision* and *Recall*

$$F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}}$$

- F_1 measure is **a harmonic mean** of the 2 metrics *Precision* and *Recall*
 - F_1 measure tends to have the value that is close to the smaller one amongst Precision and Recall
 - F_1 measure has a high value if both of Precision and Recall are high

Top-k accuracy

- Suitable for ranking (i.e., learning to rank) problems
 - E.g., Ranked lists resulted by a search/retrieval engine

$$Accuracy = \frac{1}{|D_{test}|} \sum_{x \in D_{test}} InList(ol(x), c(x))$$

- x is an example in the test set D_{test}
- $ol(x)$ is the ranked list of class labels produced by the system for the example x
- $c(x)$ is the true (expected/real) class label for the example x
- $InList(ol(x), c(x))$ is equal to 1 if the class label $c(x)$ appears in the ranked list $ol(x)$, and equal to 0 if otherwise

Nearest neighbor learning – Introduction (1)

■ Some alternative names

- Instance-based learning
- Lazy learning
- Memory-based learning

■ Nearest neighbor learning

- Given a set of training instances
 - Just store the training instances
 - Not construct a general, explicit description (model) of the target function based on the training instances
- Given a test instance (to be classified/predicted)
 - Examine the relationship between the test instance and the training ones to assign a target function value

Nearest neighbor learning – Introduction (2)

■ The input representation


- Each instance \mathbf{x} is represented as a vector in an n -dimensional vector space $\mathbf{x} \in \mathbb{R}^n$
- $\mathbf{x} = (x_1, x_2, \dots, x_n)$, where $x_i \in \mathbb{R}$ is a real number

■ We consider two learning tasks

- Nearest neighbor learning for *classification*
 - To learn a discrete-valued target function
 - The output is one of pre-defined nominal values (i.e., class labels)
- Nearest neighbor learning for *regression*
 - To learn a continuous-valued target function
 - The output is a real number

Nearest neighbor learning – Example

- 1 nearest neighbor
→ Assign z to $c2$
- 3 nearest neighbors
→ Assign z to $c1$
- 5 nearest neighbors
→ Assign z to $c1$



Nearest neighbor classifier – Algorithm

- For the classification task
- Each training instance x is represented by
 - The description: $x = (x_1, x_2, \dots, x_n)$, where $x_i \in \mathbb{R}$
 - The class label: $c \in C$, where C is a pre-defined set of class labels)
- Training phase
 - Just store the training instances set $D = \{x\}$
- Test phase. To classify a new instance z
 - For each training instance $x \in D$, compute distance between x and z
 - Compute the set $NB(z)$ – the neighbourhood of z
 - The k instances in D nearest to z according to a distance function d
 - Classify z to the majority class of the instances in $NB(z)$

Nearest neighbor regressor – Algorithm

- For the regression task (i.e., to predict a real output value)
- Each training instance x is represented by
 - The description: $x = (x_1, x_2, \dots, x_n)$, where $x_i \in \mathbb{R}$
 - The output value: $y_x \in \mathbb{R}$ (i.e., a real number)
- Training phase
 - Just store the training examples set D
- Test phase. To predict the output value for new instance z
 - For each training instance $x \in D$, compute distance between x and z
 - Compute the set $NB(z)$ – the neighbourhood of z
 - The k instances in D nearest to z according to a distance function d
 - Predict the output value of z :

$$y_z = \frac{1}{k} \sum_{x \in NB(z)} y_x$$

One vs. More than one neighbor

- Using only a single neighbor (i.e., the training instance closest to the test instance) to determine the classification/prediction is subject to errors
 - A single atypical/abnormal instance (i.e., an outlier)
 - Noise (i.e. error) in the class label (or the output value) of a single training instance
- Consider the k (>1) training instances nearest to the test one
- For a binary classification problem, the value of k is typically odd to avoid ties
 - For example, $k=3$ or $k=5$

Distance function (1)

■ The distance function d

- Play a very important role in the instance-based learning approach
- Typically defined before, and fixed through, the training and test phases – i.e., not adjusted based on data

■ Choice of the distance function d

- *Geometry* distance functions, for continuous-valued input space ($x_i \in \mathbb{R}$)
- *Hamming* distance function, for binary-valued input space ($x_i \in \{0, 1\}$)
- *Cosine similarity* function, for text classification problems (x_i is TF/IDF term weight)

Distance function (2)

■ Geometry distance functions

- Minkowski (p-norm) distance:
$$d(x, z) = \left(\sum_{i=1}^n |x_i - z_i|^p \right)^{1/p}$$

- Manhattan distance:
$$d(x, z) = \sum_{i=1}^n |x_i - z_i|$$

- Euclidean distance:
$$d(x, z) = \sqrt{\sum_{i=1}^n (x_i - z_i)^2}$$

- Chebyshev distance:
$$d(x, z) = \lim_{p \rightarrow \infty} \left(\sum_{i=1}^n |x_i - z_i|^p \right)^{1/p}$$
$$= \max_i |x_i - z_i|$$

Distance function (3)

■ *Hamming* distance function

- For binary-valued input space
- E.g., $x=(0,1,0,1,1)$

$$d(x, z) = \sum_{i=1}^n \text{Difference}(x_i, z_i)$$

$$\text{Difference}(a, b) = \begin{cases} 1, & \text{if } (a \neq b) \\ 0, & \text{if } (a = b) \end{cases}$$

■ *Cosine* similarity function

- For term weight (TF/IDF) vector

$$d(x, z) = \frac{x \cdot z}{\|x\| \|z\|} = \frac{\sum_{i=1}^n x_i z_i}{\sqrt{\sum_{i=1}^n x_i^2} \sqrt{\sum_{i=1}^n z_i^2}}$$

Attribute value normalization


- The Euclidean distance function:
$$d(x, z) = \sqrt{\sum_{i=1}^n (x_i - z_i)^2}$$
- Assume that an instance is represented by 3 attributes: Age, Income (per month), and Height (in meters)
 - $x = (\text{Age}=20, \text{Income}=12000, \text{Height}=1.68)$
 - $z = (\text{Age}=40, \text{Income}=13000, \text{Height}=1.75)$
- The distance between x and z
 - $d(x, z) = [(20-40)^2 + (12000-13000)^2 + (1.68-1.75)^2]^{1/2}$
 - The distance is dominated by the local distance (difference) on the Income attribute
 - Because the Income attribute has a large range of values
- To normalize the values of all the attributes to the same range
 - Usually the value range $[0,1]$ is used
 - E.g., for every attribute i : $x_i = x_i / \text{max_value_of_attribute_i}$

Attribute importance weight

- The Euclidean distance function: $d(x, z) = \sqrt{\sum_{i=1}^n (x_i - z_i)^2}$
 - All the attributes are considered equally important in the distance computation
- **Different attributes may have different degrees of influence on the distance metric**
- To incorporate attribute importance weights in the distance function
 - w_i is the importance weight of attribute i : $d(x, z) = \sqrt{\sum_{i=1}^n w_i (x_i - z_i)^2}$
- How to achieve the attribute importance weights?
 - By the domain-specific knowledge (e.g., indicated by experts in the problem domain)
 - By an optimization process (e.g., using a separate validation set to learn an optimal set of attribute weights)

Distance-weighted Nearest neighbor learning (1)

- Consider $NB(z)$ – the set of the k training instances nearest to the test instance z
 - Each (nearest) instance has a different distance to z
 - Should these (nearest) instances influence equally to the classification/prediction of z ? \rightarrow No!
- To weight the contribution of each of the k neighbors according to their distance to z
 - Larger weight for nearer neighbor!



Distance-weighted Nearest neighbor learning (2)

■ Let's denote v is a distance-based weighting function

- Given a distance $d(x, z)$ – the distance of x to z
- $v(x, z)$ is inversely proportional to $d(x, z)$

■ For the classification task: $$c(z) = \arg \max_{c_j \in C} \sum_{x \in NB(z)} v(x, z) \cdot \text{Identical}(c_j, c(x))$$

$$\text{Identical}(a, b) = \begin{cases} 1, & \text{if } (a = b) \\ 0, & \text{if } (a \neq b) \end{cases}$$

■ For the prediction task: $$f(z) = \frac{\sum_{x \in NB(z)} v(x, z) \cdot f(x)}{\sum_{x \in NB(z)} v(x, z)}$$

■ Select a distance-based weighting function

$$v(x, z) = \frac{1}{\alpha + d(x, z)}$$

$$v(x, z) = \frac{1}{\alpha + [d(x, z)]^2}$$

$$v(x, z) = e^{-\frac{d(x, z)^2}{\sigma^2}}$$

Lazy learning vs. Eager learning

- **Lazy learning.** The learning of the target function is postponed until the evaluation of a test (i.e., to-be-classified/predicted) example
 - To learn approximately the target function *locally* and *differently* for each to-be-classified/predicted example *at the time of the system's classification/prediction*
 - Multi times of *locally* approximate computation of the target function
 - It often takes (much) longer time to make conclusion of classification/prediction, and requires more memory resources
 - Examples: Nearest neighbor learning, Locally weighted regression
- **Eager learning.** The learning of the target function completes before the evaluation of any test (i.e., to-be classified/predicted) example
 - To learn approximately the target function *globally* for the entire examples space *at the time of the system's learning*
 - A *single and globally approximate computation* of the target function
 - Examples: Linear regression, Support vector machines, Artificial neural networks,...

Nearest neighbor learning – When?

- Examples are represented in an n -dimensional vector space \mathbb{R}^n
- The number of representation attributes is not many
- A large training set
- Advantages:
 - Very low cost for the training phase (i.e., just to store the training examples)
 - Work well for multi-label classification problems
 - Not need to learn n classifiers for n class labels
 - Nearest neighbour learning (with $k \gg 1$) can tolerate noise examples
 - Classification/prediction is done based on the k nearest neighbors
- Disadvantages:
 - To select the appropriate value of the hyper-parameter k (or appropriate dissimilarity threshold) for determining the nearest neighbors set
 - To select a distance (dissimilarity) function appropriately for a given problem
 - High computation (time, memory resource) cost at the time of the system's classification/prediction
 - May have a poor performance if irrelevant attributes are not removed