


# Cap theorem

## Scaling Traditional Databases

- Traditional RDBMSs can be either scaled:
  - o Vertically (or up).
    - Can be achieved by hardware upgrades (e.g., faster CPU, more memory, or larger disk).
    - Limited by the amount of CPU, RAM, and disk that can be configured on a single machine.
  - o Horizontally (or out).
    - Can be achieved by adding more machine.
    - Requires database sharding and probably replication.
    - Limited by the Read-to-Write ratio and communication overhead.

## Data sharding

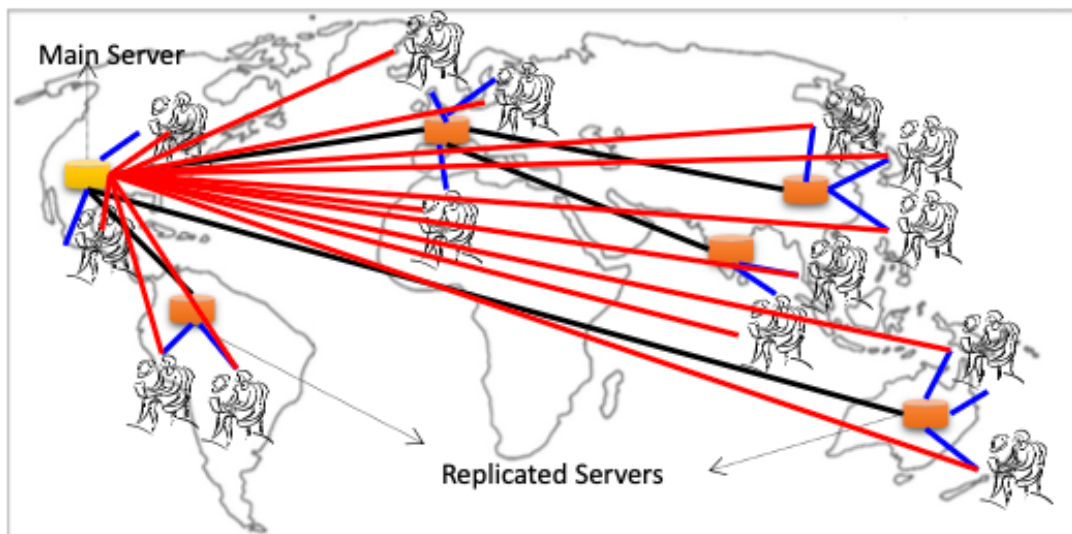
- Data is typically sharded (or striped) to allow for concurrent/parallel accesses.
- Will it scale for complex query processing?



E.g., Chunks 1, 3 and 5 can be accessed in parallel

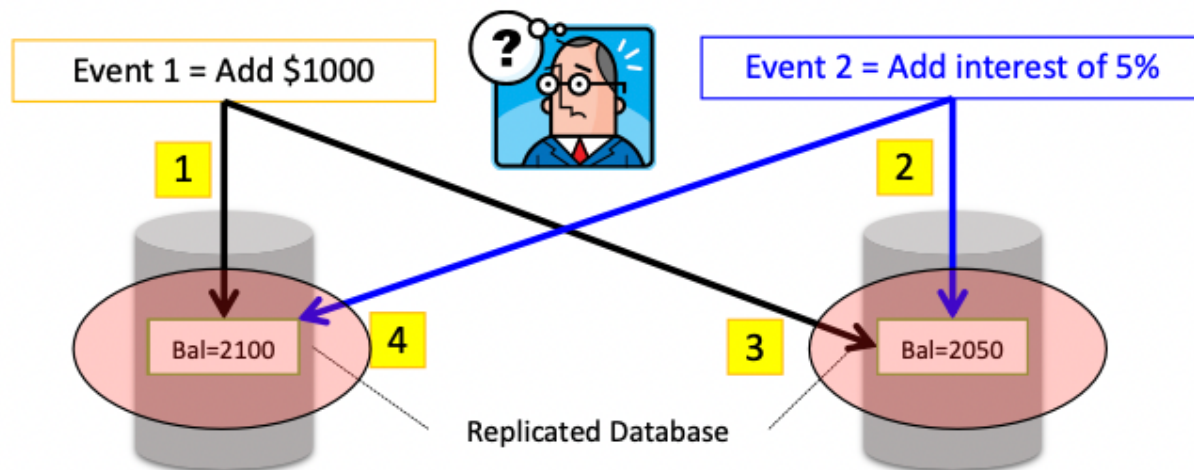
## Data replicating

- Replicating data across servers helps in:
  - o Avoiding performance bottlenecks.
  - o Avoiding single point of failures.
  - o And, hence, enhancing scalability and availability.



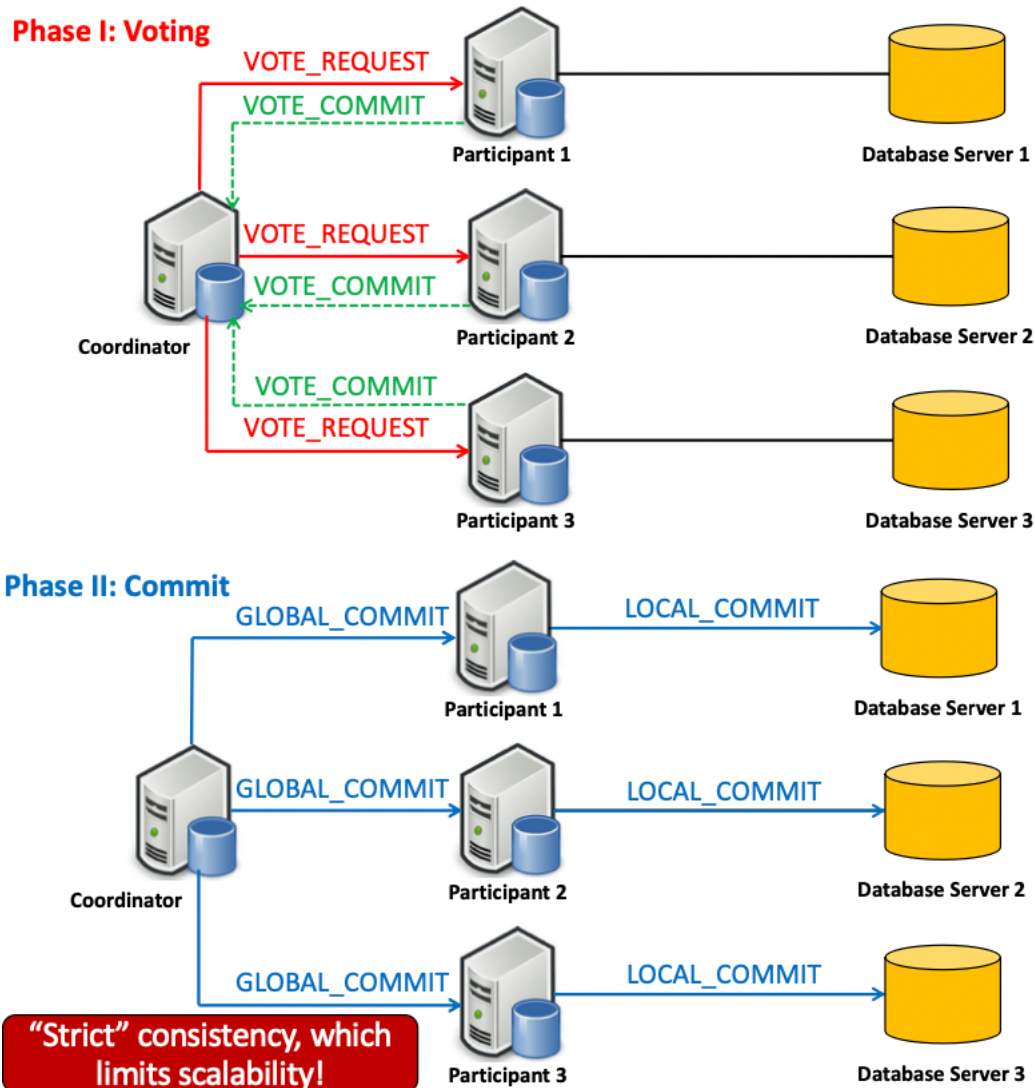
## But Consistency Becomes a Challenge

- Example:
  - In an e-commerce application, the bank database has been replicated across 2 servers.
  - Maintaining consistency of replicated data is a challenge.



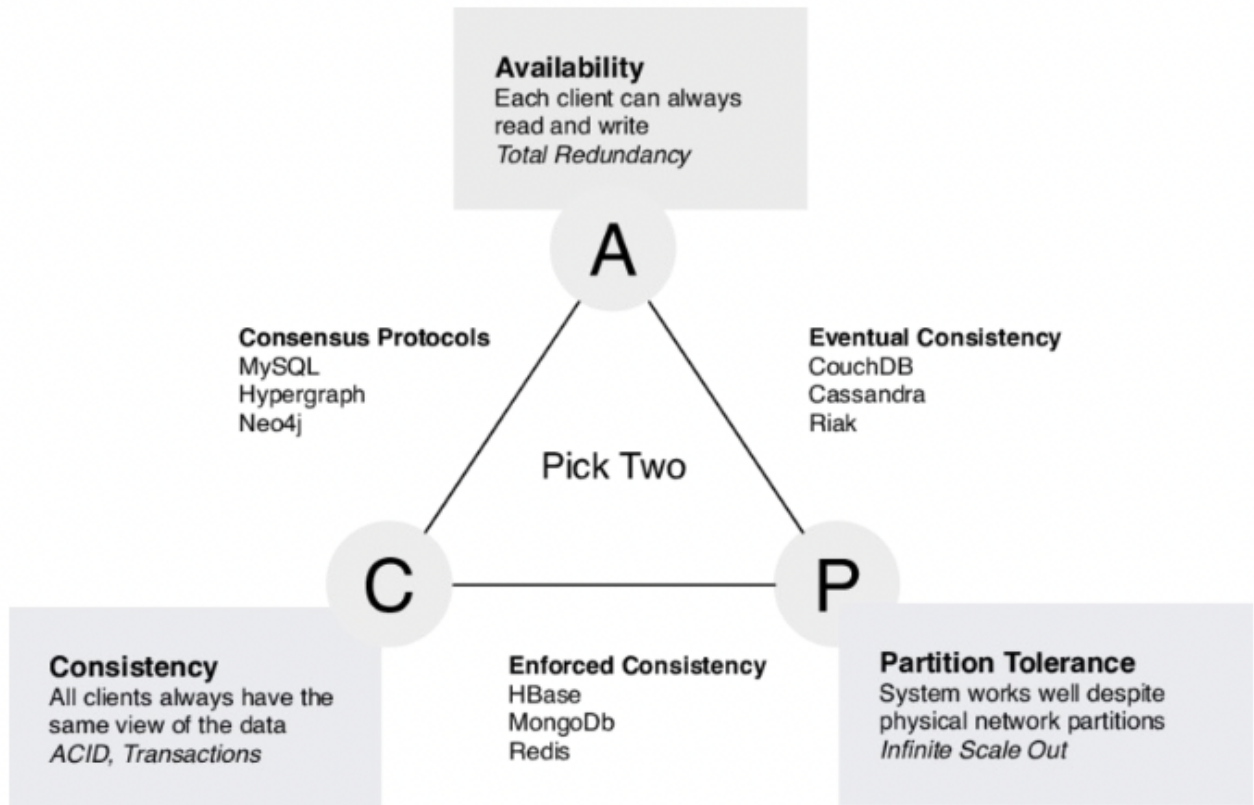
## The Two-Phase Commit Protocol

- The two-phase commit protocol (2PC) can be used to ensure atomicity and consistency.



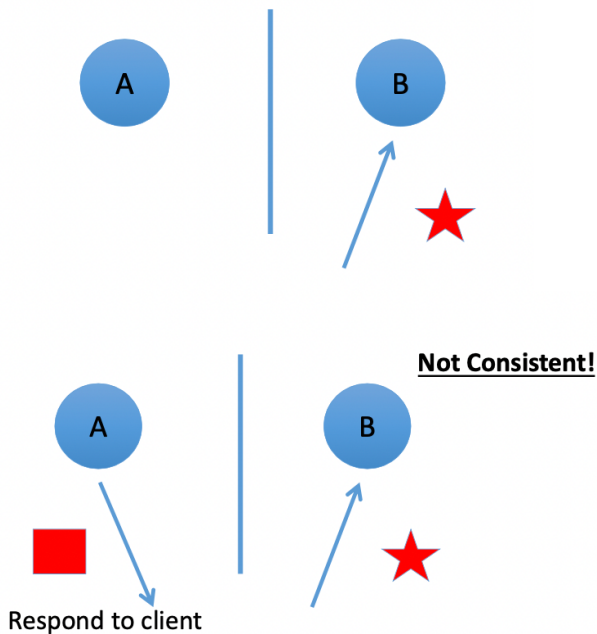
## The CAP theorem

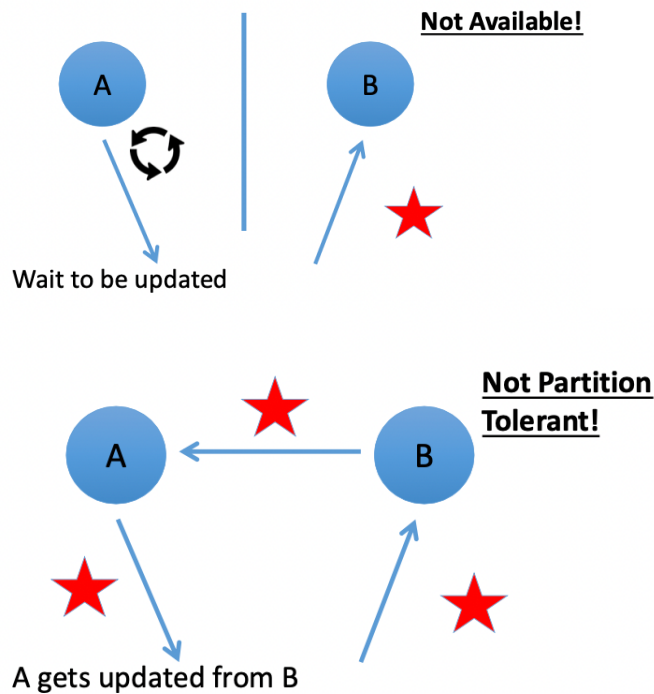
- The limitations of distributed databases can be described in the so called the CAP theorem.
  - o Consistency: every node always sees the same data at any given instance (i.e., strict consistency).
  - o Availability: the system continues to operate, even if nodes in a cluster crash, or some hardware or software parts are down due to upgrades.
  - o Partition Tolerance: the system continues to operate in the presence of network partitions.
- **CAP theorem: any distributed database with shared data can have at most two of three desirable properties, C, A, or P. These are trade-offs involved in distributed system by Eric Brewer in PODC 2000.**



## The CAP theorem: Proof

- A simple proof using two nodes:





## Scalability of relational databases

- The Relational Database is built on the principle of ACID (Atomicity, Consistency, Isolation, Durability).
- It implies that a truly distributed relational database should have **availability, consistency, and partition tolerance**.
- Which unfortunately is **impossible**.

## Large-Scale Databases

- When companies such as Google and Amazon were designing large-scale databases, 24/7 Availability was a key.
  - o A few minutes of downtime means lost revenue.
- When horizontally scaling databases to 1000s of machines, the likelihood of a node or a network failure increases tremendously.
- Therefore, to have strong guarantees on Availability and Partition Tolerance, they had to sacrifice "strict" Consistency (implied by the CAP theorem).

## Trading-Off consistency

- Maintaining consistency should balance between the strictness of consistency versus availability/scalability.
  - o Good-enough consistency depends on your application.

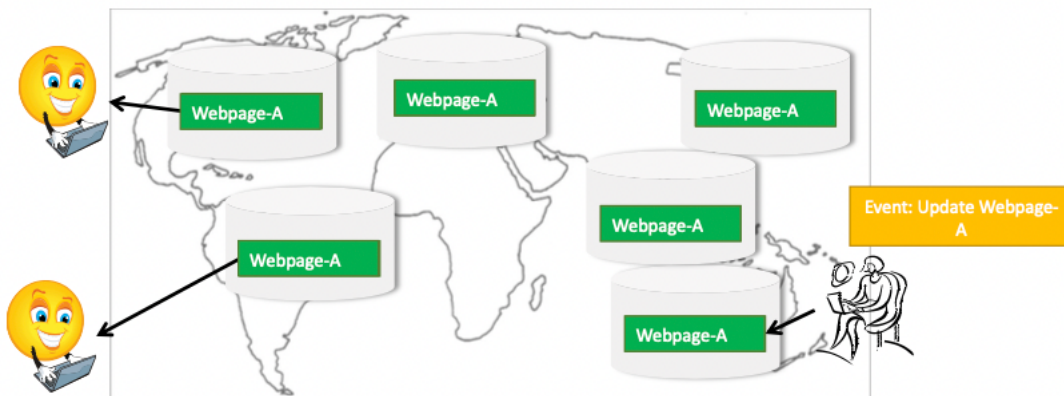


## The BASE Properties

- The CAP theorem proves that it is impossible to guarantee strict Consistency and Availability while being able to tolerate network partitions.
- This resulted in databases with relaxed ACID guarantees.
- Such databases apply the BASE properties:
  - o Basically Available: the system guarantees Availability.
  - o Soft-State: the state of the system may change over time.
  - o Eventual Consistency: the system will eventually become consistent.

## Eventually Consistency

- A database is termed as Eventually Consistent if:
  - o All replicas will gradually become consistent in the absence of new updates.



- But what if the client accesses the data from different replicas?
  - o Protocols like Read Your Own Writes (RYOW) can be applied!

## Table of Contents

<b><i>Scaling Traditional Databases .....</i></b>	<b><i>1</i></b>
<b><i>Data sharding.....</i></b>	<b><i>1</i></b>
<b><i>Data replicating.....</i></b>	<b><i>1</i></b>
<b><i>But Consistency Becomes a Challenge.....</i></b>	<b><i>2</i></b>
<b><i>The Two-Phase Commit Protocol .....</i></b>	<b><i>2</i></b>
<b><i>The CAP theorem.....</i></b>	<b><i>3</i></b>
<b><i>The CAP theorem: Proof.....</i></b>	<b><i>4</i></b>
<b><i>Scalability of relational databases .....</i></b>	<b><i>5</i></b>
<b><i>Large-Scale Databases.....</i></b>	<b><i>5</i></b>
<b><i>Trading-Off consistency.....</i></b>	<b><i>5</i></b>
<b><i>The BASE Properties .....</i></b>	<b><i>6</i></b>
<b><i>Eventually Consistency .....</i></b>	<b><i>6</i></b>