

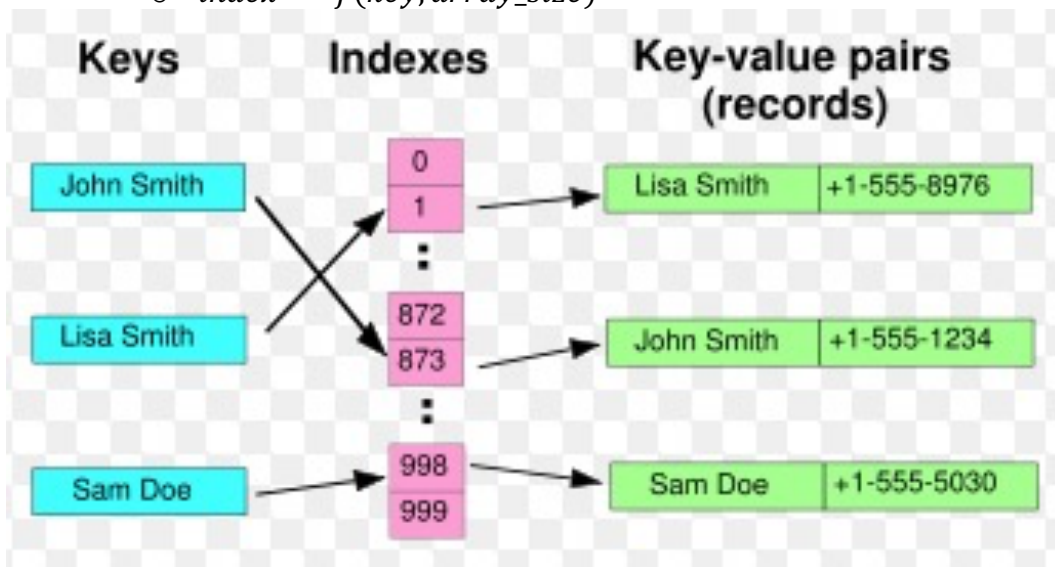
Distributed Hash Table

Outline

- Hashing.
- Distributed Hash Table.
- Chord.

A Hash Table (hash map).

- A data structure implements as an associative array that can map keys to values.
 - o Searching and insertions are $O(1)$ in the worst case.
- Uses a hash function to compute an index into an array of buckets or slots from which the correct value can be found.
 - o $index = f(key, array_size)$



Hash functions

- Crucial for good hash table performance.
- Can be difficult to achieve.
 - o WANTED: uniform distribution of hash values.
 - o A non-uniform distribution increases the number of collisions and the cost of resolving them.

Hashing for partitioning use-case


- Objective:
 - o Given document X, choose one of k servers to use.
- E.g., using modulo hashing.
 - o Number servers 1, 2, ..., k.
 - o Place X on server $i = (X \bmod k)$.

- Problem? Data may not be uniformly distributed.
 - Place X on server $i = \text{hash}(X) \bmod k$.
- Problem?
 - What happens if a server fails or joins ($k \rightarrow k \pm 1$)?
 - What If different clients have different estimate of k ?
 - Answer: All entries get remapped to new nodes!




Distributed hash table (DHT)

- Distributed Hash Table (DHT) is similar to hash table but spread across many hosts.
- Interface:
 - `insert(key, value)`
 - `lookup(key, value)`
- Every DHT node support a single operation:
 - Given key as input; route messages to node holding key.

DHT: basic idea



Neighboring nodes are “connected” at the application-level.
 Operation: take key as input; route messages to node holding key.



How to design a DHT?

- State Assignment:
 - o What (key, value) tables does a node store?
- Network Topology:
 - o How does a node select its neighbors?
- Routing Algorithm:
 - o Which neighbor to pick while routing to a destination?
- Various DHT algorithms make different choices:
 - o CAN, Chord, Pastry, Tapestry, Plaxton, Viceroy, Kademlia, Skipnet, Symphony, Koorde, Apocrypha, Land, ORDI, etc.

Chord: A scalable peer-to-peer look-up protocol for internet applications

Outline

- What is Chord?
- Consistent Hashing.
- A simple key lookup algorithm.
- Scalable key lookup algorithm.
- Node Joins and Stabilization.
- Node Fails.

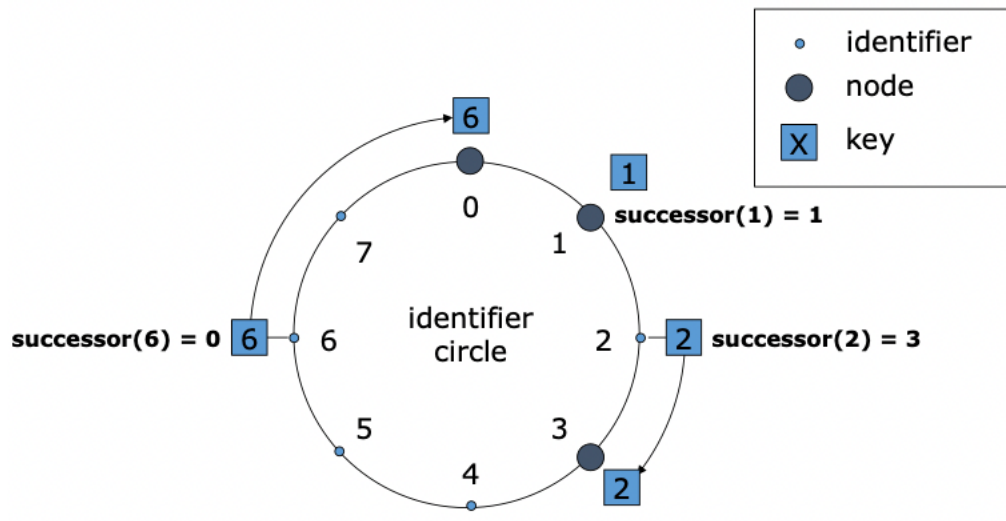
What is Chord?

- In short: a peer-to-peer lookup system.
- Given a key (data item), it maps the key onto a node (peer).
- Uses consistent hashing to assign keys to nodes.
- Solves the problem of locating key in a collection of distributed nodes.
- Maintains routing information with frequent node arrivals and departures.

Consistent hashing

- Consistent hash function assigns each node and key an m-bit identifier.
- SHA-1 is used as a base hash function.
- A node's identifier is defined by hashing the node's IP address.
- A key identifier is produced by hashing the key (chord doesn't define this. Depends on the application).
 - o $ID(\text{node}) = \text{hash}(\text{IP, Port})$.
 - o $ID(\text{key}) = \text{hash}(\text{key})$.


Consistent hashing – Successor nodes




Consistent hashing – Join and Departure

- When a node n joins the network, certain keys previously assigned to n 's successor now become assigned to n .
- When node n leaves the network, all its assigned keys are reassigned to n 's successor.

Consistent hashing – Node join



Consistent hashing – Node departure



A simple key lookup

- If each node knows only how to contact its current successor node on the identifier circle, all nodes can be visited in linear order.
- Queries for a given identifier could be passed around the circle via these successor pointers until they encounter the node that contains the key.
- Pseudo code for finding successor:

// ask node n to find the successor of id

n.find_successor(id)

if ($id \in (n, \text{successor})$)


return *successor*;

else

// forward the query around the circle

return *successor.find_successor(id)*;

- The path taken by a query from node 8 for key 54:




Scalable key location

- To accelerate lookups, Chord maintains additional routing information.
- This additional information is not essential for correctness, which is achieved as long as each node knows its correct successor.

Scalable key location – Finger tables


- Each node n' maintains a routing table with up to m entries (which is in fact the number of bits in identifiers), called finger table.
- The i^{th} entry in the table at node n contains the identity of the first node s that succeeds n by at least $2^i - 1$ on the identifier circle.
- $s = \text{successor}(n + 2^i - 1)$.
- s is called the i^{th} finger of node n , denoted by $n.\text{finger}(i)$.



- A finger table entry includes both the Chord identifier and the IP address (and port number) of the relevant node.
- The first finger of n is the immediate successor n on the circle.

Scalable key location – Example query

- The path a query for key 54 starting at node 8:



Applications: Chord-based DNS

- DNS provides a lookup service.
 - o Keys: host names, values: IP addresses.
- Chord could hash each host name to a key.
- Chord-based DNS:
 - o No special root servers.
 - o No manual management of routing information.
 - o No naming structure.
 - o Can find objects not tied to particular machines.

What is Chord? Addressed problems

- **Load balance:** chord acts as a distributed hash function, spreading keys evenly over nodes
- **Decentralization:** chord is fully distributed, no node is more important than any other, improves robustness
- **Scalability:** logarithmic growth of lookup costs with the number of nodes in the network, even very large systems are feasible
- **Availability:** chord automatically adjusts its internal tables to ensure that the node responsible for a key can always be found
- **Flexible naming:** chord places no constraints on the structure of the keys it looks up.

Summary

- Simple, powerful protocol
- Only operation: map a key to the responsible node
- Each node maintains information about $O(\log N)$ other nodes
- Lookups via $O(\log N)$ messages
- Scales well with number of nodes
- Continues to function correctly despite even major changes of the system

Table of Contents

Outline	1
A Hash Table (hash map).	1
Hash functions.....	1
Hashing for partitioning use-case	1
Distributed hash table (DHT)	2
DHT: basic idea.....	2
How to design a DHT?	4
Chord: A scalable peer-to-peer look-up protocol for internet applications.....	4
Outline	4
What is Chord?	4
Consistent hashing	4
Consistent hashing – Successor nodes.....	5
Consistent hashing – Join and Departure	5
Consistent hashing – Node join.....	5
Consistent hashing – Node departure	6
A simple key lookup.....	6
Scalable key location	7
Scalable key location – Finger tables	7
Scalable key location – Example query.....	8
Applications: Chord-based DNS.....	9
What is Chord? Addressed problems.....	9
Summary.....	9