

Spark Load Inspect

Zeppelin notebook

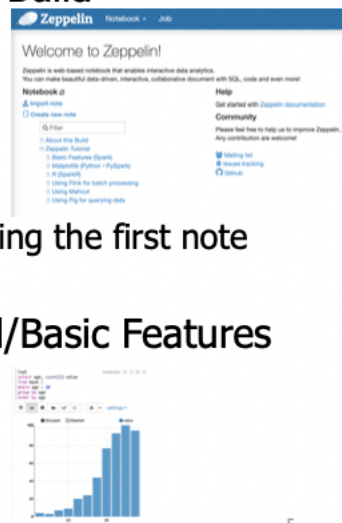
- A web-based interface for interactive data analytics.
 - o Easy to write and access your code.
 - o Support many programming languages.
 - Scala (with Apache Spark), Python (with Apache Spark), SparkSQL, Hive, Markdown, Angular, and Shell.
 - o Data visualization.
- Monitoring Spark jobs.

Installation Using Docker

- Install Docker and login
 - <https://docs.docker.com/docker-for-windows/install/>
 - <https://docs.docker.com/docker-for-mac/install/>
- Download lecture's git repository
 - <https://github.com/bk-blockchain/big-data-class>
- Run Zeppelin using docker-composer
 - docker-compose up -d --build spark_master
 - <http://localhost>

Zeppelin usage

- Run the first node: "About this Build"
 - Check Spark version
- Check Spark running mode
 - <http://localhost:4040>
 - Need to start Spark first by running the first note
- Run the second node: "Tutorial/Basic Features (Spark)"
 - Load data into table
 - SQL example



Useful Docker Commands

- **Login to a container**
 - `docker ps` (get any container id)
 - `docker exec -it container_id bash`
- **List all containers:** `docker ps -a`
- **Stop a container:** `docker stop container_id`
- **Start a stopped container:** `docker start container_id`

Load, Inspect, and Save Data

- Data is always huge that does not fit on a single machine.
 - Data is distributed on many storage nodes.
- Data scientists can likely focus on the format that their data is already in.
 - Engineers may wish to explore more output formats.
- Spark supports a wide range of input and output sources.

Data sources

- File formats and file systems.
 - Local or distributed filesystem, such as NFS, HDFS, or Amazon S3.
 - File formats including text, JSON, SequenceFiles, and protocol buffers.
- Structured data sources through Spark SQL:
 - Apache Hive.
 - Parquet.
 - JSON.
 - From RDDs.
- Databases and key/value stores.
 - Cassandra, HBase, Elasticsearch, and JDBC dbs.

File Formats

- Formats range from unstructured, like text to semi-structured, like JSON, to structured, like SequenceFiles.

Format name	Structured	Comments
Text files	No	Plain old text files. Records are assumed to be one per line.
JSON	Semi	Common text-based format, semistructured; most libraries require one record per line.
CSV	Yes	Very common text-based format, often used with spreadsheet applications.
SequenceFiles	Yes	A common Hadoop file format used for key/value data.
Protocol buffers	Yes	A fast, space-efficient multilanguage format.
Object files	Yes	Useful for saving data from a Spark job to be consumed by shared code. Breaks if you change your classes, as it relies on Java Serialization.

From Learning Spark [1]

→