

# Lecture 2: OLAP, Data Warehouse and Column Store

## Why we still study OLAP/Data Warehouse in BI?

- Understand the Big Data history.
  - o How does the requirement of (big) data analytics/business intelligence evolve over the time?
  - o What are the architecture and implementation techniques being developed? Will they still be useful in Big Data?
  - o Understand their limitation and what factors have changed from 90's to now?
- NoSQL is not only SQL.
- Hive/Impala aims to provide OLAP/BI for Big Data using Hadoop.

## Highlights

- OLAP
  - o Multi-relational Data model.
  - o Operators.
  - o SQL.
- Data warehouse (architecture, issues, optimizations).
- Join Processing.
- Column Stores (Optimized for OLAP workload).

Let's get back to the root in 70's: Relational Database

## Basic Structure

- Formally, given sets  $D_1, D_2, \dots, D_n$ . A **relation**  $r$  is a subset of  $D_1 \times D_2 \times \dots \times D_n$ .
- Thus, a **relation** is a set of n-tuples  $(a_1, a_2, \dots, a_n)$  where each  $a_i$  belong to  $D_i$ .
- Example:

```
customer_name = {Jones, Smith, Curry, Lindsay}
customer_street = {Main, North, Park}
customer_city = {Harrison, Rye, Pittsfield}
```

- Then:

```
r = {
    (Jones, Main, Harrison),
    (Smith, North, Rye),
    (Curry, North, Rye),
    (Lindsay, Park, Pittsfield)
}
```

- Is a relation over:

```
customer_name, customer_street, customer_city
```

## Relation Schema

- $A_1, A_2, \dots, A_n$  are **attributes**.
- $R = (A_1, A_2, \dots, A_n)$  is a **relation schema**.  
Example:
  - o  $\text{customer\_schema} = (\text{customer\_name}, \text{customer\_street}, \text{customer\_city})$
- $r(R)$  is a relation on the relation schema  $R$ .  
Example:
  - o  $\text{customer}(\text{customer\_schema})$

## Relation Instance

- The current values (**relation instance**) of a relation are specified by a table.
- An element  $t$  of  $r$  is a **tuple**, represented by a row in a table.

The diagram shows a table representing a relation instance. The table has three columns labeled *customer\_name*, *customer\_street*, and *customer\_city*. The rows contain the following data:

<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>
Jones	Main	Harrison
Smith	North	Rye
Curry	North	Rye
Lindsay	Park	Pittsfield

Annotations point to the table structure:

- An arrow points from the text "attributes (or columns)" to the column headers.
- An arrow points from the text "tuples (or rows)" to the first four rows of the table.
- The word "customer" is centered below the table.

## Database

- A **database** consists of multiple relations.
- Information about an enterprise is broken up into parts, with each relation storing one part of the information.

```
account : stores information about accounts
depositor : stores information about which customer owns which account
customer : stores information about customers
```
- Storing all information as a single relation such as:  
**bank(account\_number, balance, customer\_name, ...)**  
results in **repetition** of information (e.g., two customers own an account) and the need for null values (e.g., represent a customer without an account).

## Banking Example

```
branch (branch-name, branch-city, assets)
customer (customer-name, customer-street, customer-city)
account (account-number, branch-name, balance)
loan (loan-number, branch-name, amount)
depositor (customer-name, account-number)
borrower (customer-name, loan-number)
```

## Relational Algebra

- Primitives:
  - Projection ( $\pi$ )
  - Selection ( $\sigma$ )
  - Cartesian product ( $\times$ )
  - Set union ( $\cup$ )
  - Set difference ( $-$ )
  - Rename ( $\rho$ )
- Other operations:
  - Join ( $\bowtie$ )
  - Group by... aggregation
  - ...

## What happens next?

- SQL.
- System R (DB2), INGRES, ORACLE, SQL-Server, Teradata.
  - B+-Tree (select).
  - Transaction Management.
  - Join Algorithm.

## In early 90's: OLAP & Data Warehouse

## Database Workloads

- OLTP (online transaction processing).
  - Typical applications: e-commerce, banking, airline reservations.
  - User-facing: real-time, low latency, **highly concurrent**.
  - Tasks: relatively small set of "standard" transactional queries.
  - Data access pattern: random reads, updates, writes (involving relatively small amounts of data).
- OLAP (online analytical processing).
  - Typical applications: business intelligence, data mining.
  - Back-end processing: **batch workloads, less concurrency**.
  - Tasks: complex analytical queries, often ad-hoc.
  - Data access pattern: table scans, large amounts of data involved per query.

## OLTP

- Most database operations involve (On-Line Transaction Processing).
  - o Short, simple, frequent queries and/or modifications, each involving a small number of tuples.
  - o Examples: Answering queries from Web interface, sales at cash registers, selling airline tickets.

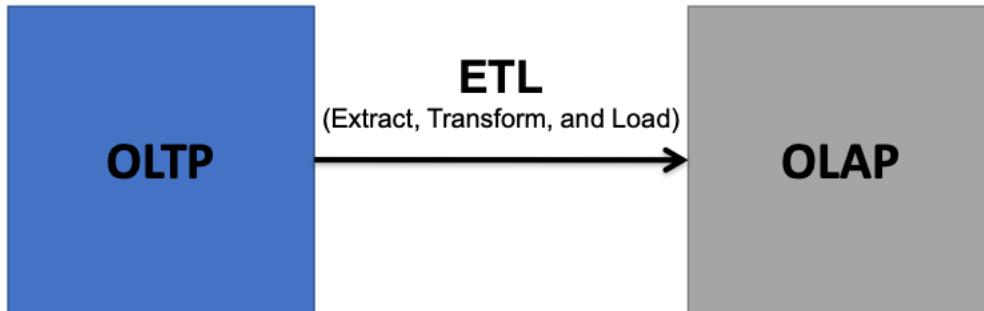
## OLAP

- Of increasing importance are On-line Application Processing (OLAP) queries.
  - o Few, but complex queries – may run for hours.
  - o Queries do not depend on having an **ABSOLUTELY** up-to-date database.
- Examples:
  - o 1. Amazon analyzes purchases by its customers to come up with an individual screen with products of likely interest to the customer.
  - o 2. Analysts at Wal-Mart look for items with increasing sales in some region.

## One Database or Two?

- Downsides of co-existing OLTP and OLAP workloads:
  - o Poor memory management.
  - o Conflicting data access patterns.
  - o Variable latency.
- Solution: separate databases.
  - o User-facing OLTP database for high-volume transaction.
  - o Data warehouse for OLAP workloads
  - o How do we connect the two?

## OLTP/OLAP Architecture



## OLTP/OLAP Integration

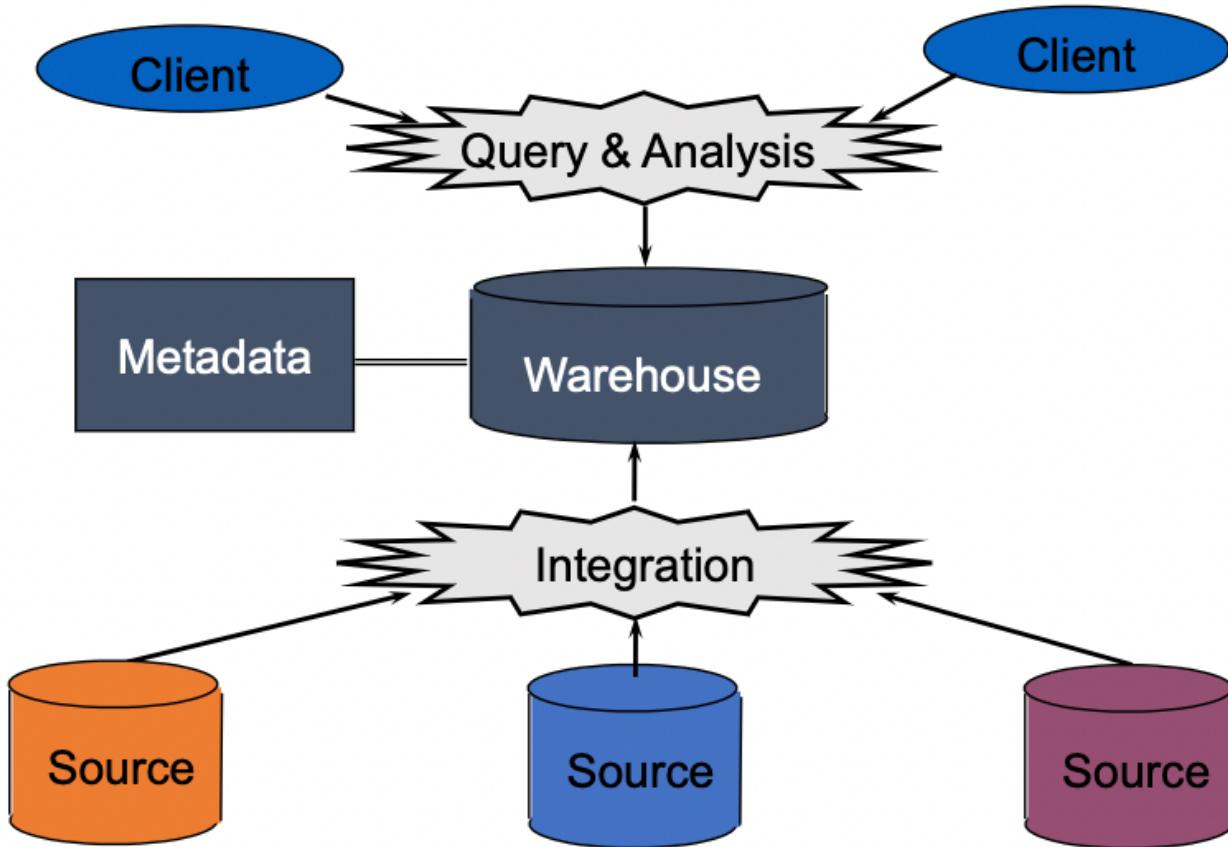
- OLTP database for user-facing transactions.
  - o Retain records of all activity.
  - o Periodic ETL (e.g., nightly).
- Extract-Transform-Load (ETL):
  - o Extract records from source.
  - o Transform: clean data, check integrity, aggregate, etc.
  - o Load into OLAP database.
- OLAP database for data warehousing:

- BI: reporting, ad-hoc queries, data mining, etc.
- Feedback to improve OLTP services.

## The Data Warehouse

- The most common form of data integration.
  - Copy sources into a single DB (warehouse) and try to keep it up to date.
  - Usual method: periodic reconstruction of the warehouse, perhaps overnight.
  - Frequently essential for analytic queries.

## Warehouse Architecture



## Star Schemas

- A **star schema** is a common organization for data at a warehouse. It consists of:
  - 1. Fact table:** a very large accumulation of facts such as sales.
    - Often “insert-only”.
  - 2. Dimension tables:** smaller, generally static information about the entities involved in the facts.

## Example: Star Schema

- Suppose we want to record in a warehouse information about every beer sale: the bar, the brand of beer, the drinker who bought the beer, the day, the time, and the price charged.
- The fact table is a relation:

**Sales(bar, beer, drinker, day, time, price)**

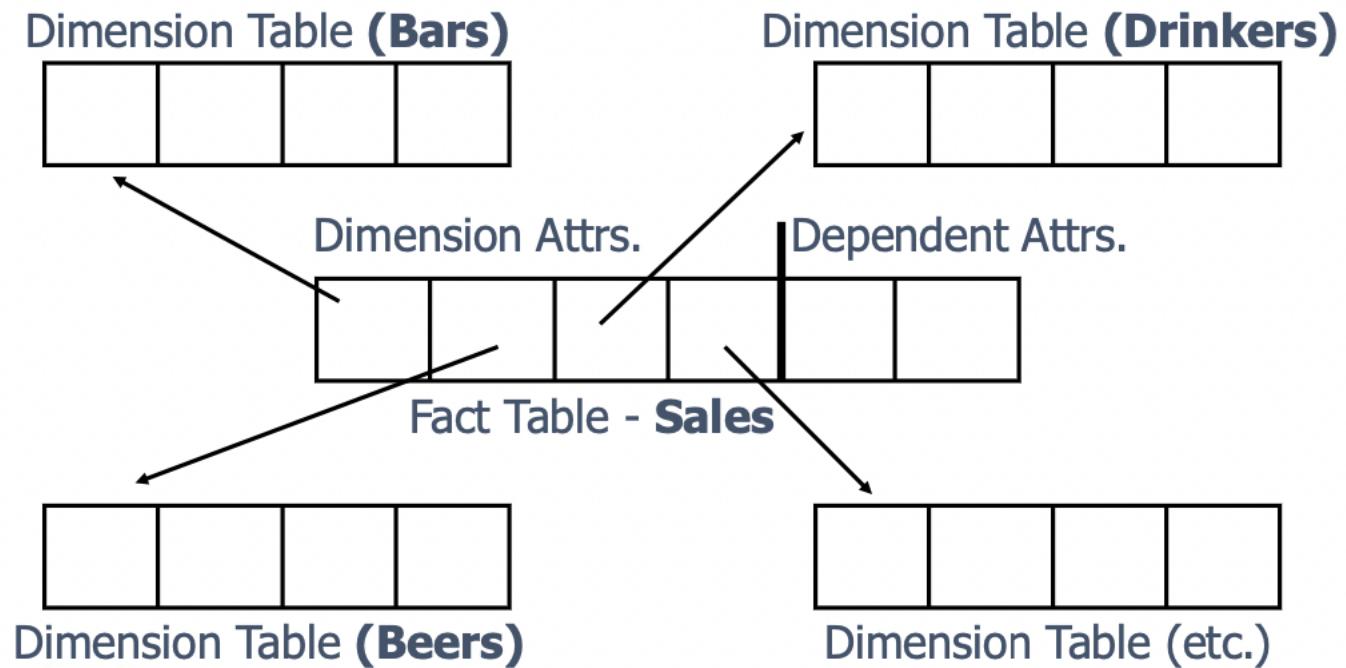
- The dimension tables include information about the bar, beer, and drinker “dimensions”:

**Bars(bar, addr, license)**

**Beers(beer, manf)**

**Drinkers(drinker, addr, phone)**

Visualization – Start Schema



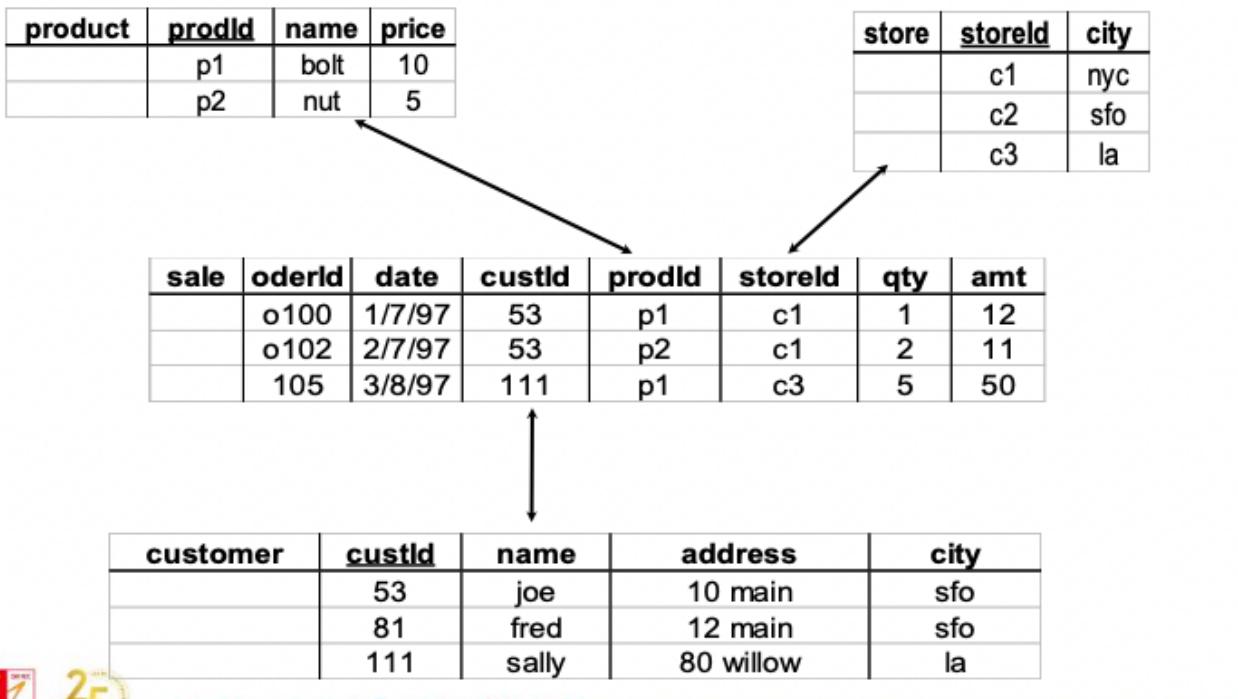
### Dimension and Dependent Attributes

- Two classes of fact-table attributes:
  - o **1. Dimension attributes:** the keys of dimension tables.
  - o **2. Dependent attributes:** a value determined by the dimension attributes of tuple.

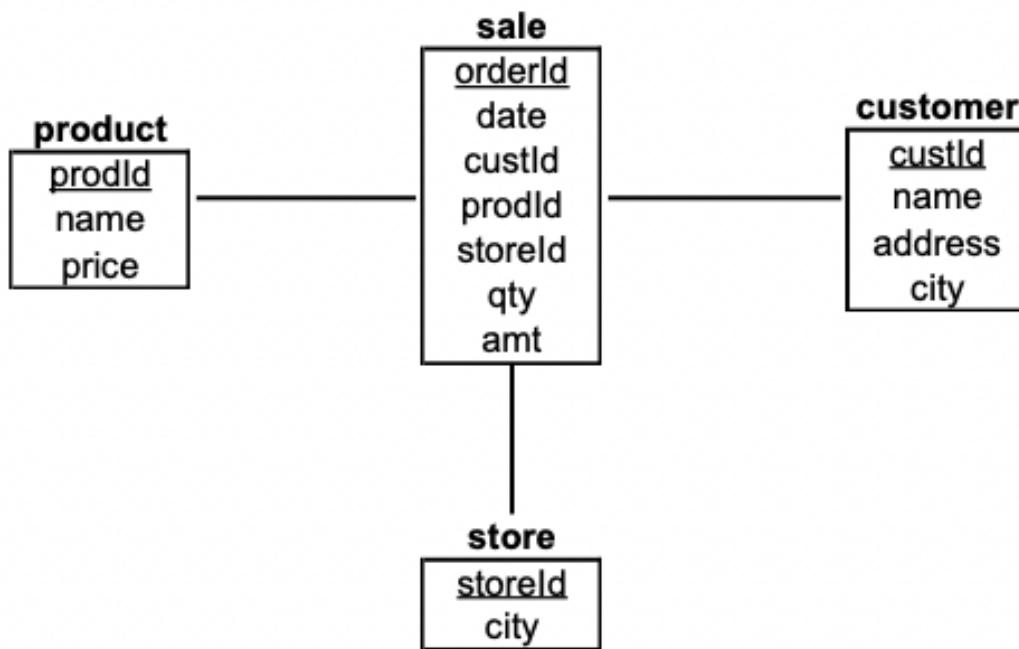
### Warehouse Models & Operators

- Data Models:
  - o Relations.
  - o Stars & snowflakes.
  - o Cubes.
- Operators:
  - o Slice & dice.
  - o Roll-up, drill-down.
  - o Pivoting.
  - o Other.

## Star



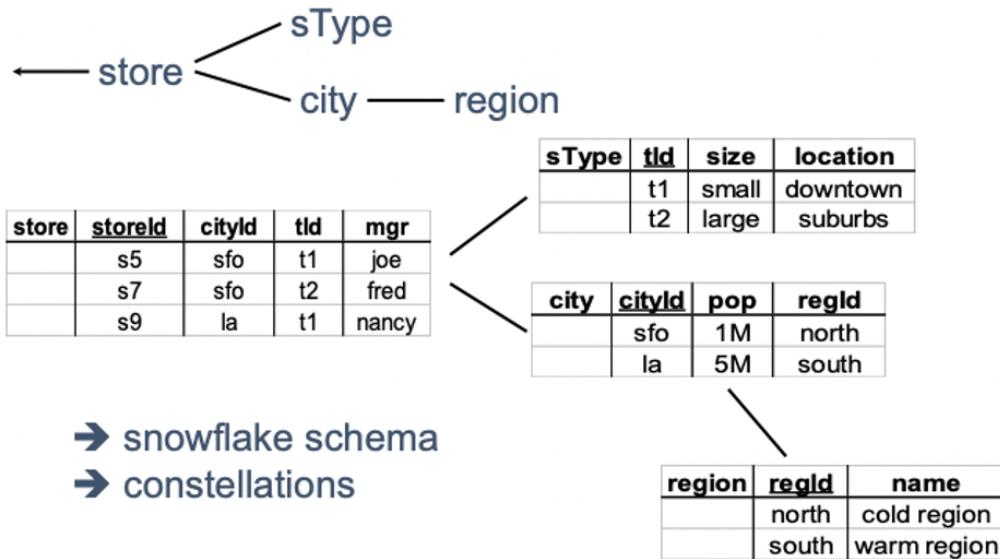
## Star Schema



## Terms

- Fact table.
- Dimension tables.
- Measures.

## Dimension Hierarchies



## Aggregates

- In SQL: **SELECT sum(amt) FROM SALE WHERE date = 1**

sale	prodId	storeId	date	amt
	p1	c1	1	12
	p2	c1	1	11
	p1	c3	1	50
	p2	c2	1	8
	p1	c1	2	44
	p1	c2	2	4

→ 81

**SELECT date, sum(amt) FROM SALE GROUP BY date**

sale	prodId	storeId	date	amt
	p1	c1	1	12
	p2	c1	1	11
	p1	c3	1	50
	p2	c2	1	8
	p1	c1	2	44
	p1	c2	2	4

→

ans	date	sum
	1	81
	2	48

**SQL: SELECT date, sum(amt) FROM SALE GROUP BY date, prodId**

sale	prodId	storeId	date	amt
	p1	c1	1	12
	p2	c1	1	11
	p1	c3	1	50
	p2	c2	1	8
	p1	c1	2	44
	p1	c2	2	4

→

sale	prodId	date	amt
	p1	1	62
	p2	1	19
	p1	2	48

→ rollup →

← drill-down ←

## ROLAP VS. MOLAP

- ROLAP: Relational On-Line Analytical Processing
- MOLAP: Multi-Dimensional On-Line Analytical Processing

### Cube

Fact table view:

sale	prodId	storeId	amt
	p1	c1	12
	p2	c1	11
	p1	c3	50
	p2	c2	8

Multi-dimensional cube:

	c1	c2	c3
p1	12		50
p2	11	8	

dimensions = 2

### 3-D Cube

Fact table view:

sale	prodId	storeId	date	amt
	p1	c1	1	12
	p2	c1	1	11
	p1	c3	1	50
	p2	c2	1	8
	p1	c1	2	44
	p1	c2	2	4

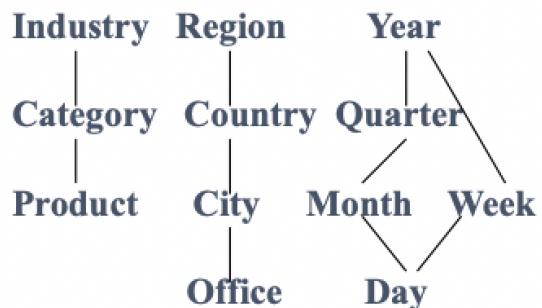
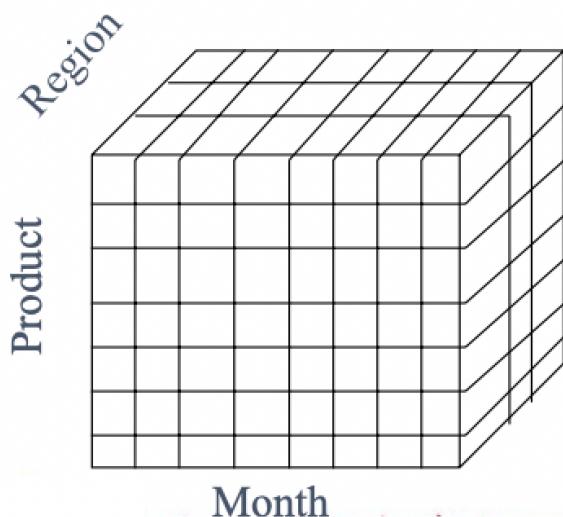
Multi-dimensional cube:

day 2		c1	c2	c3
	p1	44	4	
day 1		c1	c2	c3
	p1	12		50
	p2	11	8	

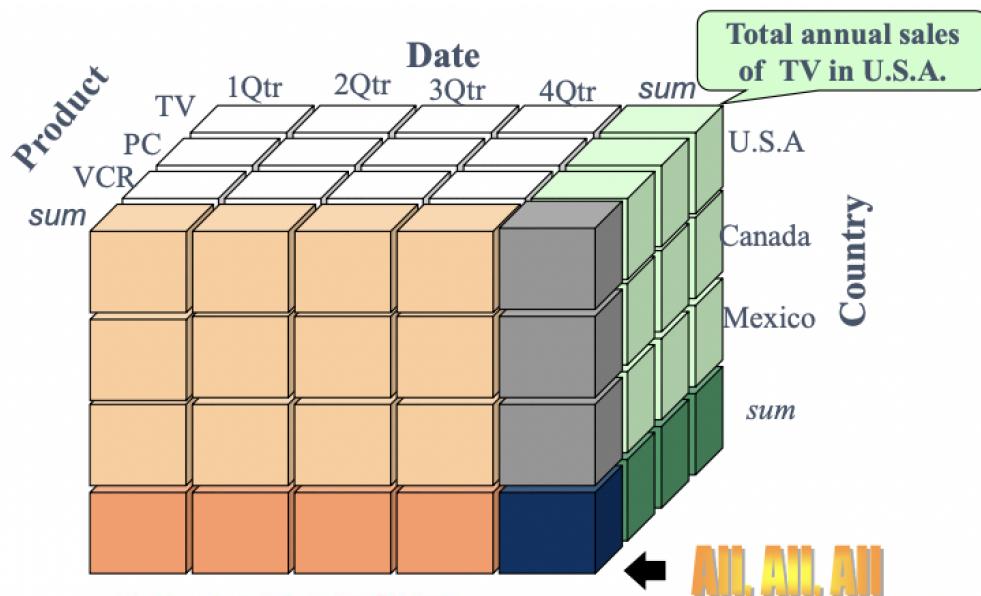
### Multidimensional Data

- Sales volume as a function of product, month, and region.

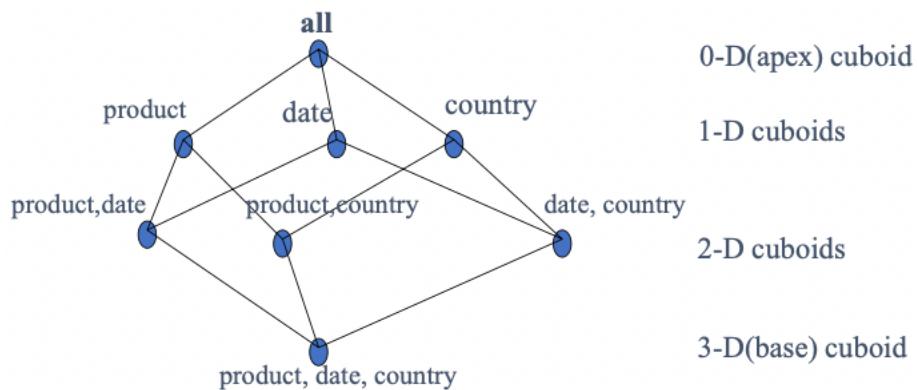
Dimensions: Product, Location, Time  
Hierarchical summarization paths



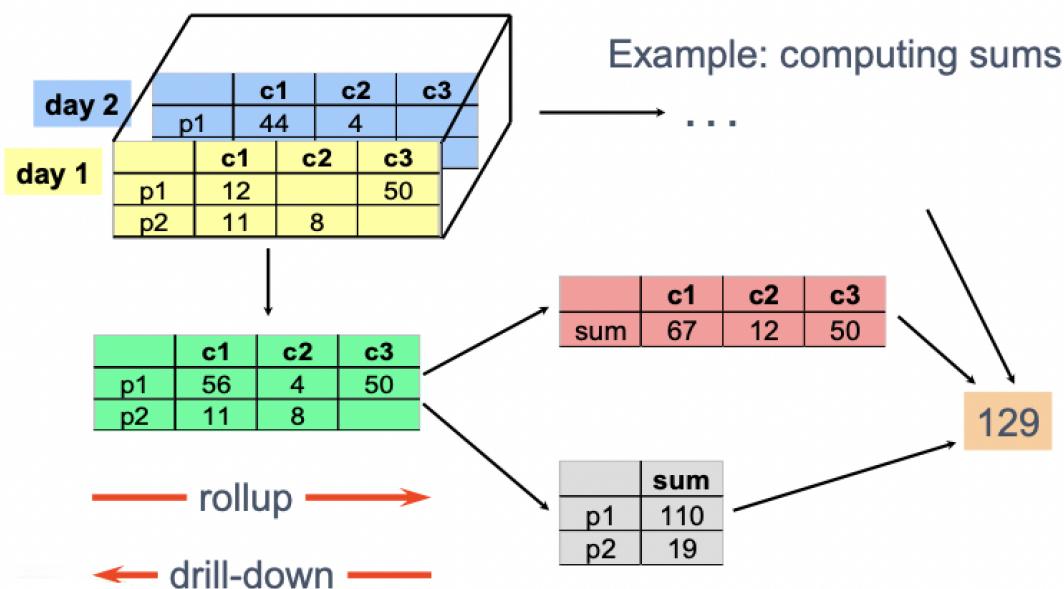
## A Sample Data Cube



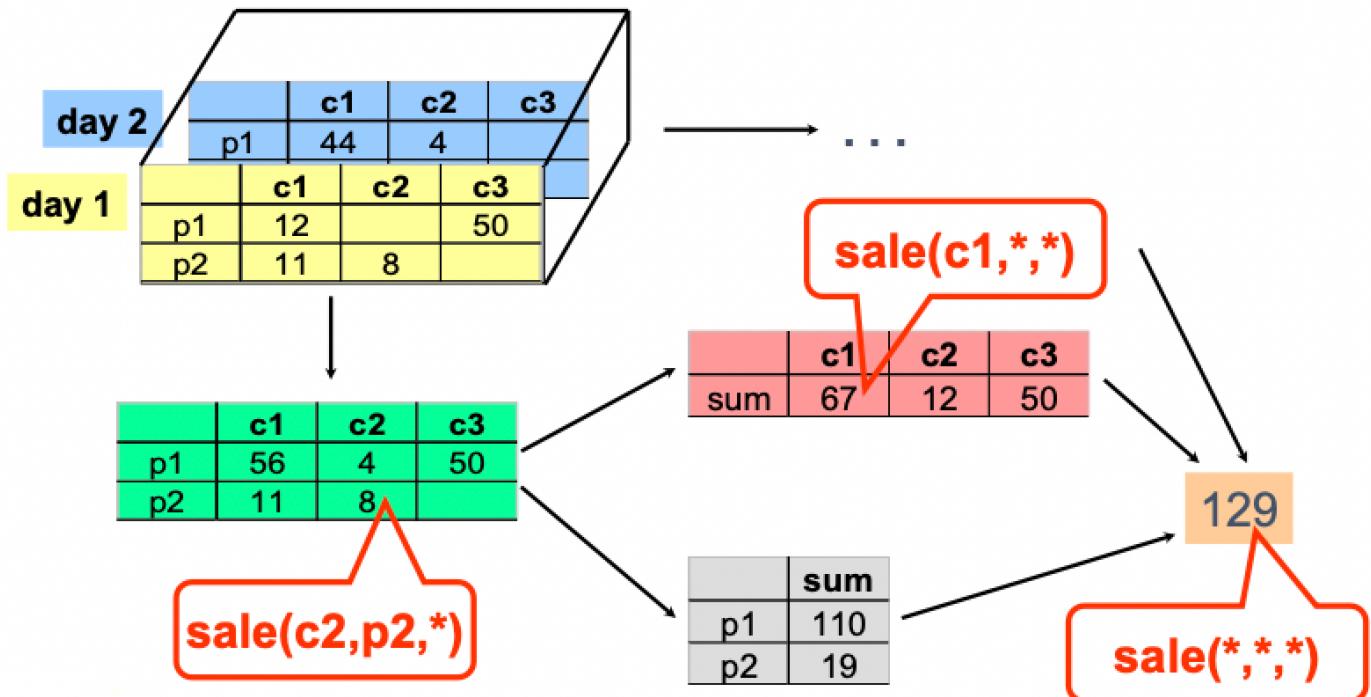
## Cuboids Corresponding to the Cube



## Cube Aggregation



## Cube Operators

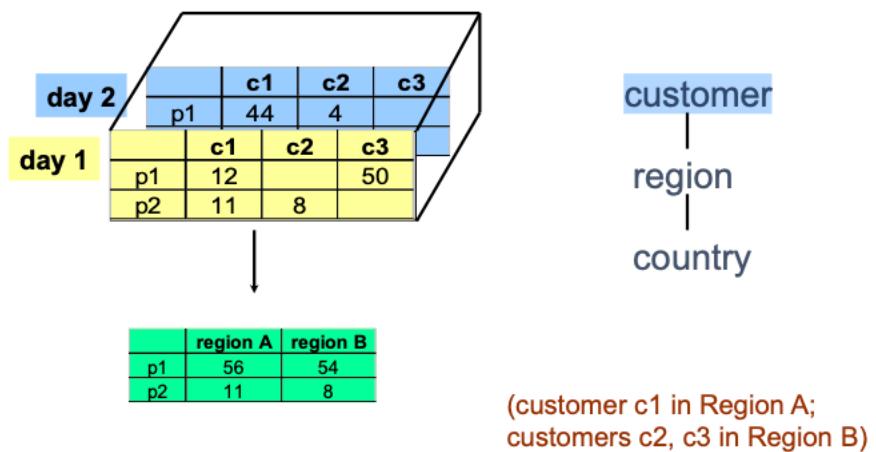


## Extended Cube

	*	c1	c2	c3	*	sum
day 2		p1	56	4	50	110
		p2	11	8		19
day 1		p1	44	4		48
		p2	11	8		19
*		*	23	8	50	81

**sale(\*, p2, \*)**

## Aggregation Using Hierarchies

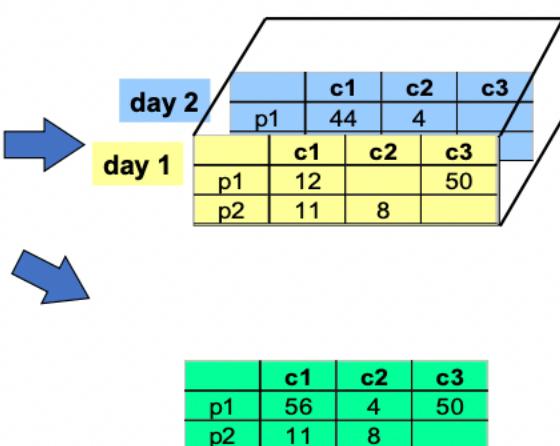


## Pivoting

Fact table view:

sale	prodId	storeId	date	amt
	p1	c1	1	12
	p2	c1	1	11
	p1	c3	1	50
	p2	c2	1	8
	p1	c1	2	44
	p1	c2	2	4

Multi-dimensional cube:



CUBE Operator (SQL-99)

Chevy Sales Cross Tab				
Chevy	1990	1991	1992	Total (ALL)
black	50	85	154	289
white	40	115	199	354
Total (ALL)	90	200	353	1286

```

SELECT model, year, color, sum(sales) as sales
FROM   sales
WHERE  model in ( 'Chevy' )
AND    year BETWEEN 1990 AND 1992
GROUP BY CUBE (model, year, color);
  
```

- Computes union of 8 different groups:
  - o {(model, year, color), (model, year), (model, color), (year, color), (model), (year), (color), ()}

## Aggregates

- Operators: sum, count, max, min, median, average.
- “Having” clause.
- Cube (& Rollup) operator.
- Using dimension hierarchy.
  - average by region (within store).
  - maximum by month (within date).

## Query & Analysis Tools

- Query Building.
- Report Writers (comparisons, growth, graphs, etc.)

- Spreadsheet Systems.
- Web Interfaces.
- Data Mining.

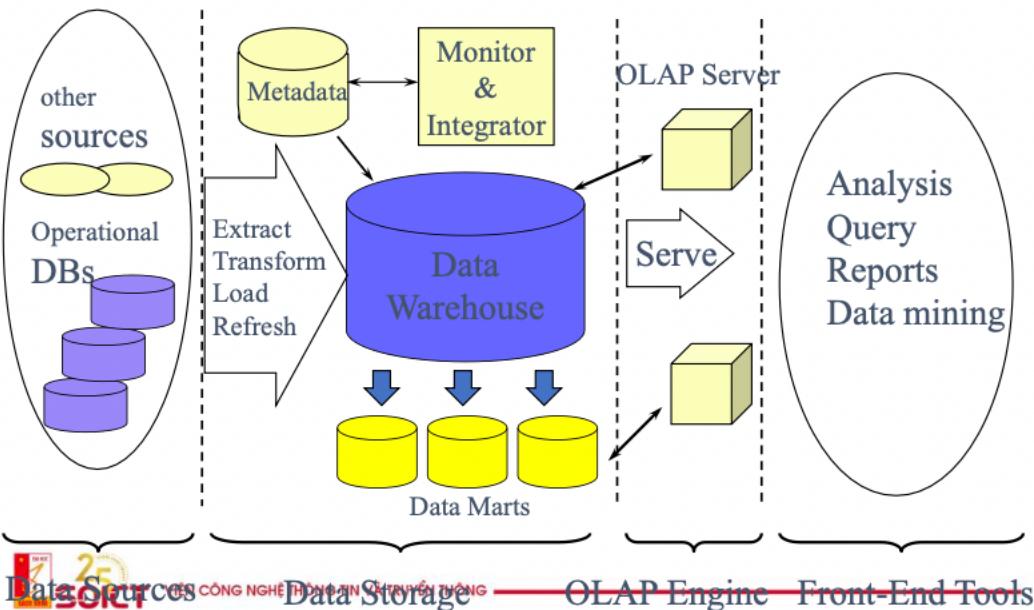
## Other Operations

- Time functions.
  - o E.g., time average.
- Computed Attributes:
  - o E.g., commission = sales \* rate.
- Text Queries:
  - o E.g., find documents with words X AND B.
  - o E.g., rank documents by frequency of words X, Y, Z.

## Data Warehouse Implementation

- Monitoring: Sending data from sources.
- Integrating: Loading, Cleansing.
- Processing: Query processing, indexing, etc.
- Managing: Metadata, Design.

## Multi-Tiered Architecture



## Monitoring

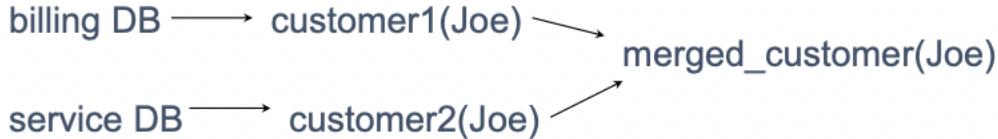
- Source Types: relational, flat file, IMS, VSAM, IDMS, WWW, newswire, etc.
- Incremental vs. Refresh.

customer	id	name	address	city
	53	joe	10 main	sfo
	81	fred	12 main	sfo
	111	sally	80 willow	la



## Data Cleansing

- Migration (e.g., yen → dollars).
- Scrubbing: use domain-specific knowledge (e.g., social security numbers).
- Fusion (e.g., mail list, customer merging).
- Auditing: discover rules & relationships (like data mining).



## Loading data

- Incremental vs. refresh.
- Off-line vs. on-line.
- Frequency of loading.
- At night, 1x a week/month, continuously.
- Parallel/Partitioned load.

## OLAP Implementation

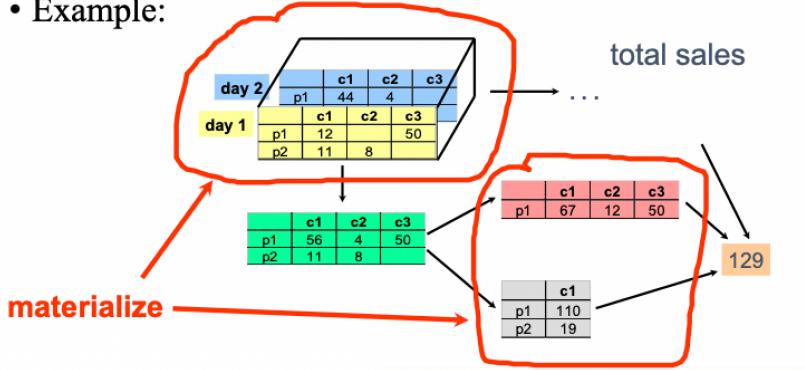
### Derived Data

- Derived Warehouse Data
  - o Indexes
  - o Aggregates
  - o materialized views
- When to update derived data?
- Incremental vs. refresh.

### What to Materialize

- Store in warehouse results useful for common queries.

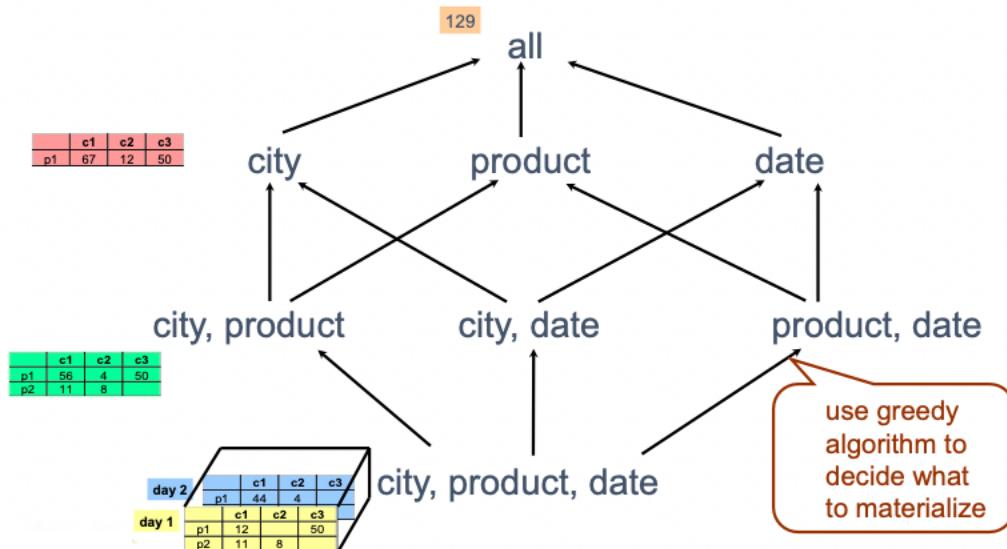
• Example:



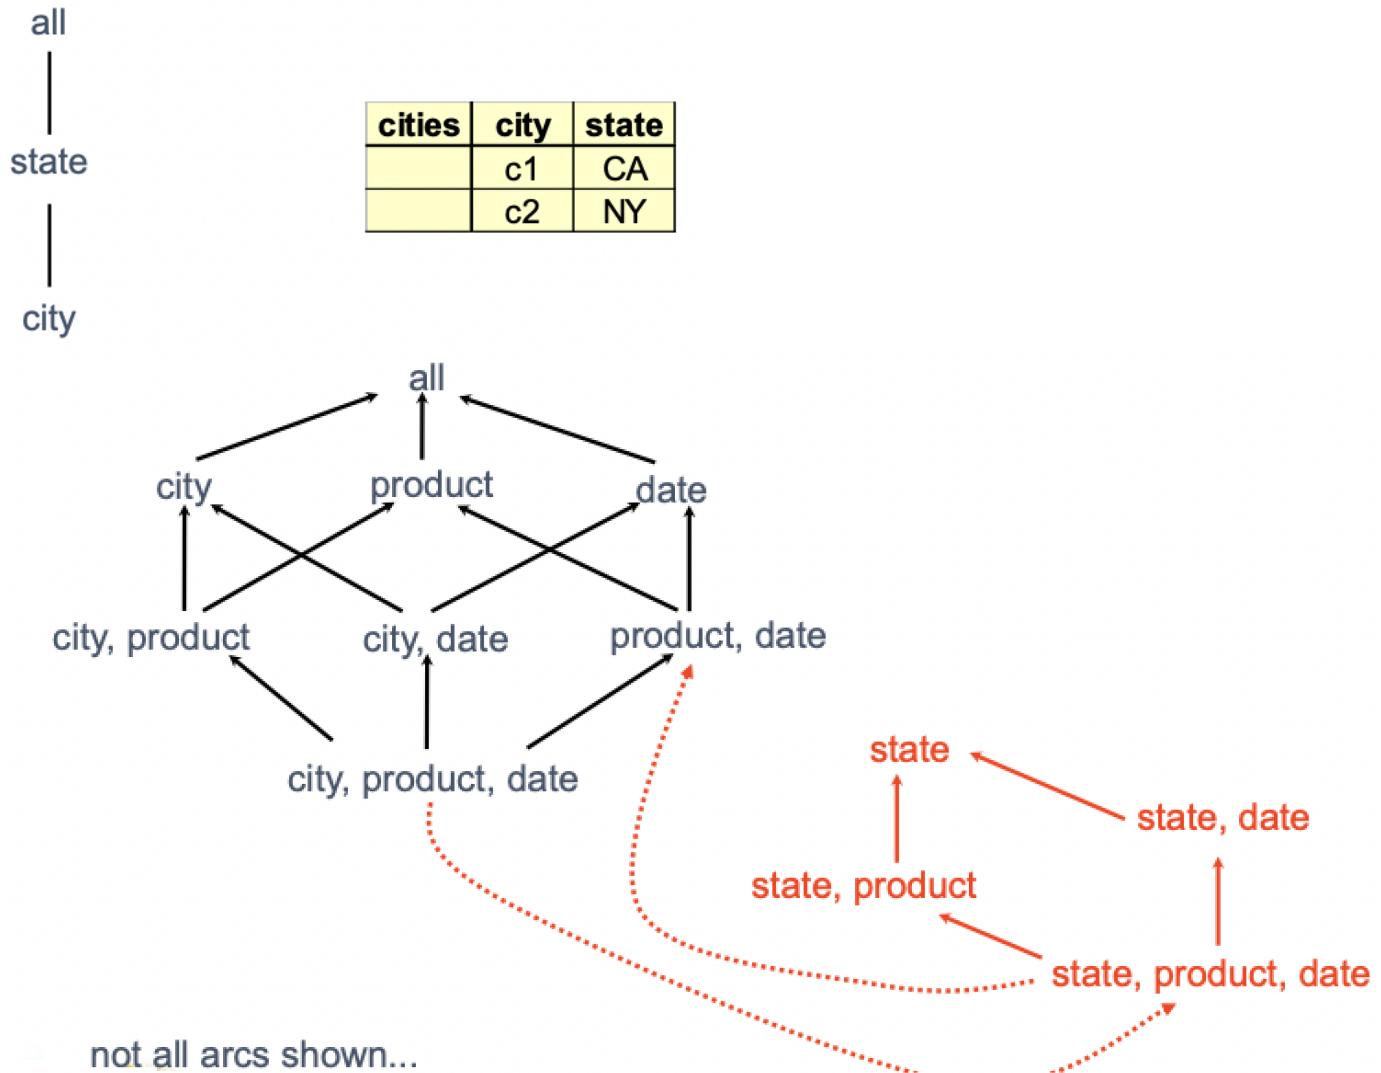
### Materialization Factors

- Type/frequency of queries.
- Query response time.
- Storage Cost.
- Update Cost.

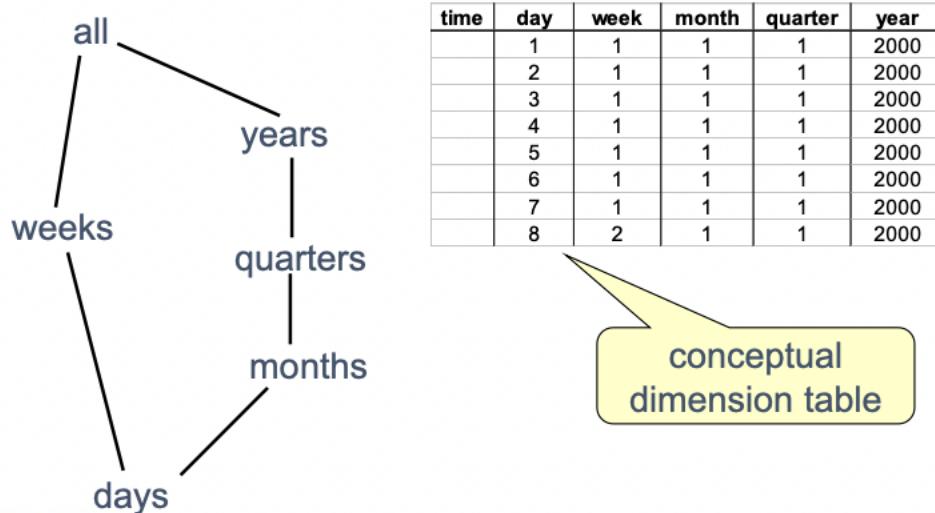
## Cube Aggregates Lattice



## Dimension Hierarchies



## Interesting Hierarchy



## Indexing OLAP Data: Bitmap Index

- Index on a particular column
- Each value in the column has a bit vector: bit-op is fast
- The length of the bit vector: # of records in the base table
- The i-th bit is set if the i-th row of the base table has the value for the indexed column
- not suitable for high cardinality domains

Base table			Index on Region			Index on Type			
Cust	Region	Type	RecID	Asia	Europe	America	RecID	Retail	Dealer
C1	Asia	Retail	1	1	0	0	1	1	0
C2	Europe	Dealer	2	0	1	0	2	0	1
C3	Asia	Dealer	3	1	0	0	3	0	1
C4	America	Retail	4	0	0	1	4	1	0
C5	Europe	Dealer	5	0	1	0	5	0	1

## Join Processing

- How does DBMS join 2 tables?
- Sorting is one way.
- Database must choose the best way for each query.

## Schema for Examples

Sailors (sid: integer, sname: string, rating: integer, age: real)  
 Reserves (sid: integer, bid: integer, day: dates, rname: string)

- Like old schema, `rname` added for variations.
- Reserves:
  - o Each tuple is 40 bytes long.
  - o 100 tuples per page.
  - o M = 1000 pages total.
- Sailors:

- Each tuple is 50 bytes long.
- 80 tuples per page.
- $N = 500$  pages total.

## Equality Joins with One Join Column

```
SELECT *
FROM   Reserves R1, Sailors S1
WHERE  R1.sid=S1.sid
```

- In algebra:  $R \bowtie S$ . Common! Must be carefully optimized.  
 $R \times S$  is large; so,  $R \times S$  followed by a selection is inefficient.
- Assume:  $M$  tuples in  $R$ ,  $p_R$  tuples per page,  $N$  tuples in  $S$ ,  $p_S$  tuples per page.
  - In our examples,  $R$  is Reserves and  $S$  is Sailors.
- We will consider more complex join conditions later.
- **Cost metric:** # of I/Os. We will ignore output costs.

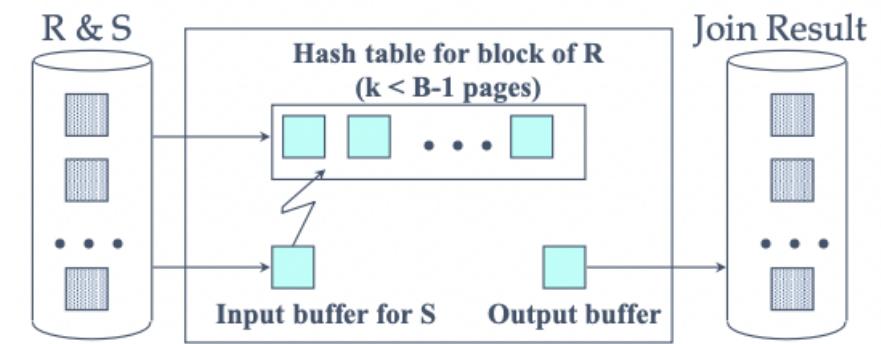
## Simple Nested Loops Join

```
foreach tuple r in R do
  foreach tuple s in S do
    if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```

- For each tuple in the *outer* relation  $R$ , we scan the entire *inner* relation  $S$ .
  - Cost:  $M + p_R * M * N = 1000 + 100 * 1000 * 500$  I/Os.
- Page-oriented Nested Loops join: For each *page* of  $R$ , get each *page* of  $S$ , and write out matching pairs of tuples  $\langle r, s \rangle$ , where  $r$  is in  $R$ -page and  $S$  is in  $S$ -page.
  - Cost:  $M + M * N = 1000 + 1000 * 500$
  - If smaller relation ( $S$ ) is outer, cost =  $500 + 500 * 1000$

## Block Nested Loops Join

- Use one page as an input buffer for scanning the inner  $S$ , one page as the output buffer, and use all remaining pages to hold ‘‘block’’ of outer  $R$ .
- For each matching tuple  $r$  in  $R$ -Block,  $s$  in  $S$ -page, add  $\langle r, s \rangle$  to result. Then read next  $R$ -block, scan  $S$ , etc.



## Examples of Block Nested Loops

- Cost: Scan of outer + #outer blocks \* scan of inner
  - #outer blocks =  $\lceil \# \text{ of pages of outer} / \text{blocksize} \rceil$
- With Reserves (R) as outer, and 100 pages of R:
  - Cost of scanning R is 1000 I/Os; a total of 10 blocks.
  - Per block of R, we scan Sailors (S); 10\*500 I/Os.
  - If space for just 90 pages of R, we would scan S 12 times.
- With 100-page block of Sailors as outer:
  - Cost of scanning S is 500 I/Os; a total of 5 blocks.
  - Per block of S, we scan Reserves; 5\*1000 I/Os.
- With sequential reads considered, analysis changes: may be best to divide buffers evenly between R and S.

## Index Nested Loops Join

```
foreach tuple r in R do
    foreach tuple s in S where  $r_i == s_j$  do
        add  $\langle r, s \rangle$  to result
```

- If there is an index on the join column of one relation (say S), can make it the inner and exploit the index.
  - Cost:  $M + (M * p_R) * \text{cost of finding matching S tuples}$
- For each R tuple, cost of probing S index is about 1.2 for hash index, 2-4 for B+ tree. Cost of then finding S tuples (assuming Alt. (2) or (3) for data entries) depends on clustering.
  - Clustered index: 1 I/O (typical), unclustered: upto 1 I/O per matching S tuple.

## Examples of Index Nested Loops

- Hash-index (Alt. 2) on *sid* of Sailors (as inner):
  - Scan Reserves: 1000 page I/Os, 100\*1000 tuples.
  - For each Reserves tuple: 1.2 I/Os to get data entry in index, plus 1 I/O to get (the exactly one) matching Sailors tuple.  
Total: 220,000 I/Os.
- Hash-index (Alt. 2) on *sid* of Reserves (as inner):
  - Scan Sailors: 500 page I/Os, 80\*500 tuples.
  - For each Sailors tuple: 1.2 I/Os to find index page with data entries, plus cost of retrieving matching Reserves tuples.  
Assuming uniform distribution, 2.5 reservations per sailor (100,000 / 40,000). Cost of retrieving them is 1 or 2.5 I/Os depending on whether the index is clustered.

## Sort-Merge Join

$$(R \bowtie S)_{i=j}$$

- Sort R and S on the join column, then scan them to do a ``merge`` (on join column), and input result tuples.
  - o Advance scan of R until current R-tuple  $\geq$  current S tuple, then advance scan of S until current S-tuple  $\geq$  current R tuple; do this until current R tuple = current S tuple.
  - o At this point, all R tuples with same value in  $R_i$  (current R group) and all S tuples with same value in  $S_j$  (current S group) match; output  $\langle r, s \rangle$  for all pairs of such tuples.
  - o • Then resume scanning R and S.
- R is scanned once; each S group is scanned once per matching R tuple. (Multiple scans of an S group are likely to find needed pages in buffer).

### Example of Sort-Merge Join

		<u>sid</u>	<u>bid</u>	<u>day</u>	<u>rname</u>
<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>		
22	dustin	7	45.0	28	103
28	yuppy	9	35.0	28	103
31	lubber	8	55.5	31	101
44	guppy	5	35.0	31	102
58	rusty	10	35.0	31	101
				58	103
					11/12/96
					dustin

- **Cost:**  $M \log M + N \log N + (M+N)$ 
  - The cost of scanning,  $M+N$ , could be  $M*N$  (very unlikely!)
- With 35, 100 or 300 buffer pages, both Reserves and Sailors can be sorted in 2 passes; total join cost: 7500.

### Refinement of Sort-Merge Join

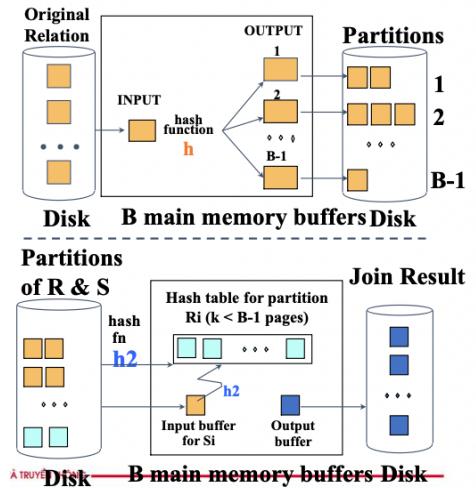
- We can combine the merging phases in the *sorting* of R and S with the merging required for the join.
  - With  $B > \sqrt{L}$ , where  $L$  is the size of the larger relation, using the sorting refinement that produces runs of length  $2B$  in Pass 0, #runs of each relation is  $< B/2$ .
  - Allocate 1 page per run of each relation, and 'merge' while checking the join condition.
  - **Cost:** read+write each relation in Pass 0 + read each relation in (only) merging pass (+ writing of result tuples).
  - In example, cost goes down from 7500 to 4500 I/Os.
- In practice, cost of sort-merge join, like the cost of external sorting, is *linear*.

## Hash-Join

- Partition both relations using hash fn  $h$ : R tuples in partition i will only match S tuples in partition i.
- Read in a partition of R, hash it using  $h_2 (< h!)$ . Scan matching partition of S, search for matches.

## Observations in Hash-Join

- #partitions  $k < B-1$  (why?), and  $B-2 > \text{size of largest partition}$  to be held in memory. Assuming uniformly sized partitions, and maximizing k, we get:
  - $k = B-1$ , and  $M/(B-1) < B-2$ , i.e., B must be  $> \sqrt{M}$
- If we build an in-memory hash table to speed up the matching of tuples, a little more memory is needed.
- If the hash function does not partition uniformly, one or more R partitions may not fit in memory. Can apply hash-join technique recursively to do the join of this R-partition with corresponding S-partition.

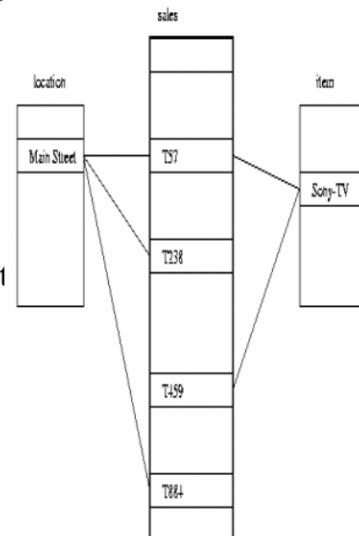


## Cost of Hash-Join

- In partitioning phase, read+write both relns;  $2(M+N)$ . In matching phase, read both relns;  $M+N$  I/Os.
- In our running example, this is a total of 4500 I/Os.
- Sort-Merge Join vs. Hash Join:
  - Given a minimum amount of memory (*what is this, for each?*) both have a cost of  $3(M+N)$  I/Os. Hash Join superior on this count if relation sizes differ greatly. Also, Hash Join shown to be highly parallelizable.
  - Sort-Merge less sensitive to data skew; result is sorted.

## Join Indices

- Traditional indices map the values to a list of record ids
  - It materializes relational join in JI file and speeds up relational join — a rather costly operation
- In data warehouses, join index relates the values of the dimensions of a star schema to rows in the fact table.
  - E.g. fact table: *Sales* and two dimensions *city* and *product*
    - A join index on *city* maintains for each distinct city a list of R-IDs of the tuples recording the Sales in the city
  - Join indices can span multiple dimensions



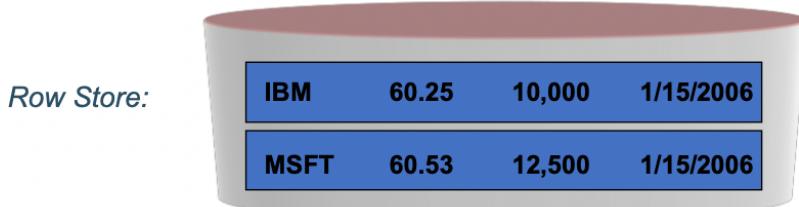
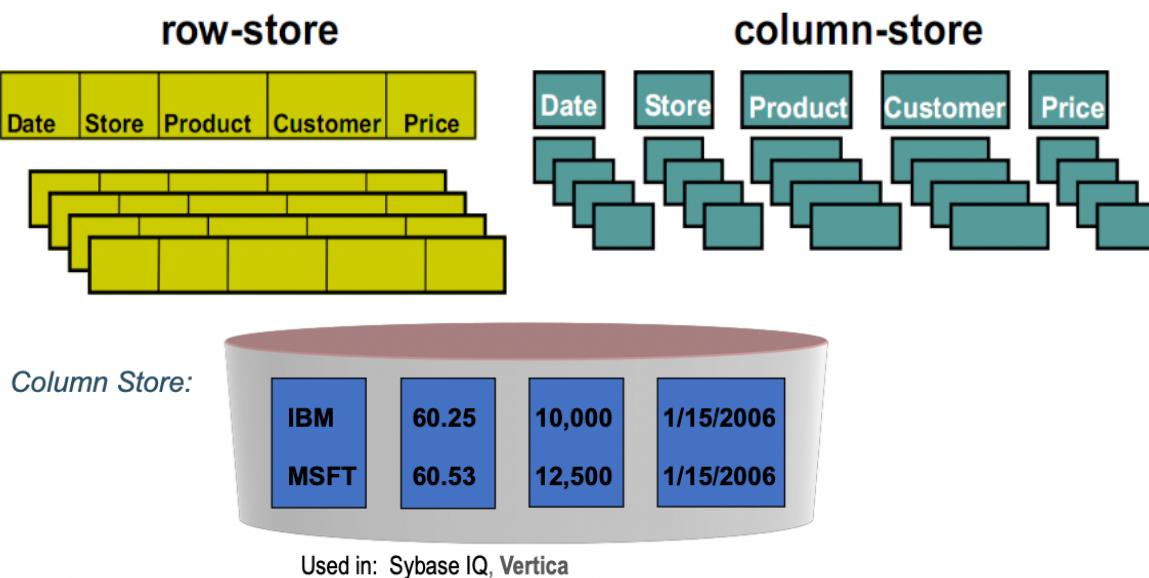
## General Join Conditions

- Equalities over several attributes (e.g.,  $R.sid=S.sid$  AND  $R.rname=S.sname$ ):
  - For Index NL, build index on  $\langle sid, sname \rangle$  (if S is inner); or use existing indexes on  $sid$  or  $sname$ .
  - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns.
- Inequality conditions (e.g.,  $R.rname < S.sname$ ):
  - For Index NL, need (clustered!) B+ tree index.
    - Range probes on inner; # matches likely to be much higher than for equality joins.
  - Hash Join, Sort Merge Join not applicable.
  - Block NL quite likely to be the best join method here.

An invention in 2000s: Column Stores for OLAP

## Row Store and Column Store

- In row store data are stored in the disk tuple by tuple.
- Where in column store data are stored in the disk column by column



For example the query

```
SELECT account.account_number,
       sum(usage.toll_airtime),
       sum(usage.toll_price)
  FROM usage, toll, source, account
 WHERE usage.toll_id = toll.toll_id
   AND usage.source_id = source.source_id
   AND usage.account_id = account.account_id
   AND toll.type_ind in ('AE', 'AA')
   AND usage.toll_price > 0
   AND source.type != 'CIBER'
   AND toll.rating_method = 'IS'
   AND usage.invoice_date = 20051013
 GROUP BY account.account_number
```

Row-store: one row = 212 columns!

Column-store: 7 attributes

Row Store	Column Store
(+) Easy to add/modify a record	(+) Only need to read in relevant data
(-) Might read in unnecessary data	(-) Tuple writes require multiple accesses

- So, column storages are suitable for read-mostly, read intensive, large data repositories.

## Column Stores: High Level

- Read only what you need:
  - o “Fat” fact tables are typical.
  - o Analytics read only a few columns.
- Better compression.
- Execute and compressed data.
- Materialized views help row stores and column stores about equally.

## Data Model (Vertica/C-Store)

- Same as relational data model:
  - o Tables, rows, columns.
  - o Primary keys and foreign keys.
  - o **Projections:**
    - From single table.
    - Multiple joined tables.
- Example:

## Possible C-store model

### Normal relational model

**EMP**(name, age, dept, salary)

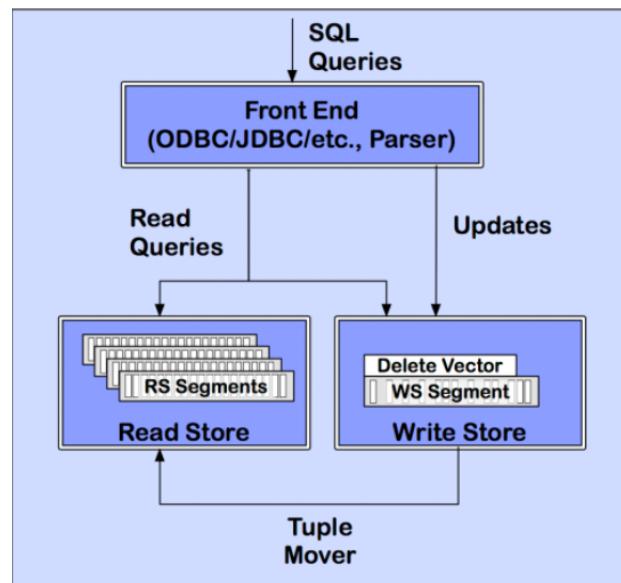
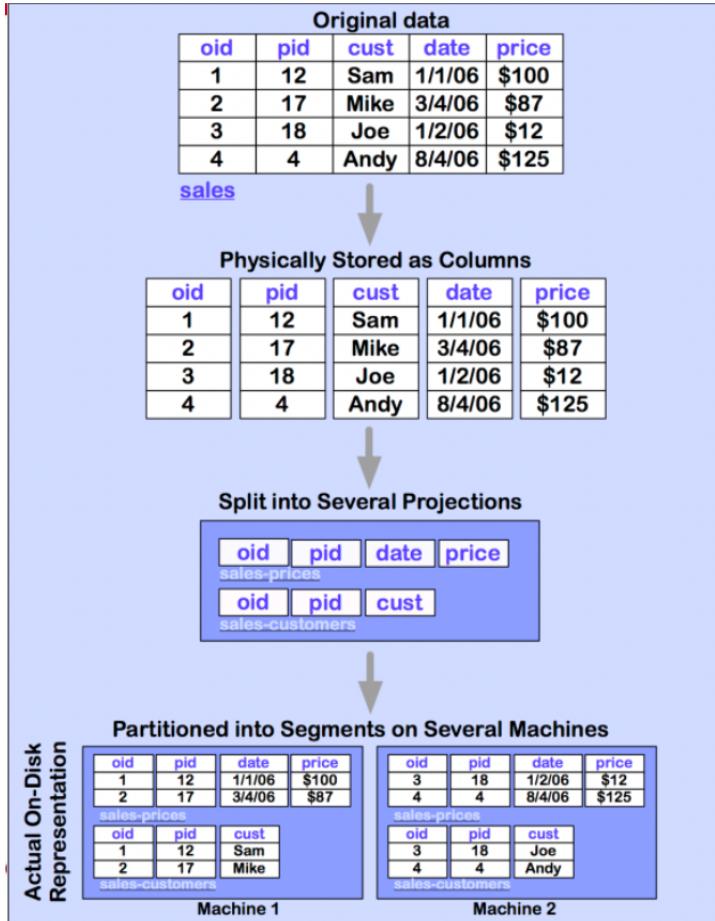
**DEPT**(dname, floor)

**EMP1** (name, age)

**EMP2** (dept, age, DEPT.floor)

**EMP3** (name, salary)

**DEPT1**(dname, floor)



### Read store: Column Encoding/Compression

- Use compression schemes and indices:
  - o Null Suppression.
  - o Dictionary encoding.
  - o Run Length encoding.
  - o Bit-Vector encoding.
  - o Self-order (key), few distinct values
    - (value, position, # items)
    - Indexed by clustered B-tree
  - o Foreign-order (non-key), few distinct values
    - (value, bitmap index)
    - B-tree index: position → values
  - o Self-order, many distinct values
    - Delta from the previous value
    - B-tree index
  - o Foreign-order, many distinct values
    - Unencoded

## Compression

- Trades I/O for CPU
  - o Increased column-store opportunities.
  - o Higher data value locality in column stores.
  - o Techniques such as run length encoding far more useful.

## Write Store

- Same structure, but explicitly use (segment, key) to identify records.
  - o Easier to maintain the mapping.
  - o Only concerns the inserted records.
- Tuple mover:
  - o Copies batch of records to RS.
- Delete record:
  - o Mark it on RS.
  - o Purged by tuple mover.

## How to solve read/write conflict

- Situation: one transaction updates the record X, while another transaction reads X.
- Use snapshot isolation.

## Query Execution – Operators

- Select: Same as relational algebra but produces a bit string.
- Project: Same as relational algebra.
- Join: Joins projections according to predicates.
- Aggregation: SQL like aggregates.
- Sort: Sort all columns of a projection.
- Decompress: Converts compressed column to uncompressed representation.
- Mask(Bitstring B, Projection Cs) => emit only those values whose corresponding bits are 1.
- Concat: Combines one or more projections sorted in the same order into a single projection.
- Permute: Permutes a projection according to the ordering defined by a join index.
- Bitstring operators: Band – Bitwise AND, Bor – Bitwise OR, Bnot – complement.

## Benefits in Query Processing

- Selection – has more indices to use.
- Projection – some “projections” already defined.
- Join – some projections are materialized joins.
- Aggregations – works on required columns only.

## Evaluation

- Use TPC-H – decision support queries.
- Storage

C-Store	Row Store	Column Store
1.987 GB	4.480 GB	2.650 GB

## Query Performance

Query	C-Store	Row Store	Column Store
Q1	0.03	6.80	2.24
Q2	0.36	1.09	0.83
Q3	4.90	93.26	29.54
Q4	2.09	722.90	22.23
Q5	0.31	116.56	0.93
Q6	8.50	652.90	32.83
Q7	2.54	265.80	33.24

- Row store uses materialized views

Query	C-Store	Row Store	Column Store
Q1	0.03	0.22	2.34
Q2	0.36	0.81	0.83
Q3	4.90	49.38	29.10
Q4	2.09	21.76	22.23
Q5	0.31	0.70	0.63
Q6	8.50	47.38	25.46
Q7	2.54	18.47	6.28

Summary: The performance gain

- Column representation: avoid reads of unused attributes.
- Storing overlapping projections – multiple orderings of a column, more choices for query optimization.
- Compression of data – more orderings of a column in the same amount of space.
- Query operators operate on compressed representation.

## Google's Dremel/Big Query: Interactive Analysis of Web-Scale Datasets

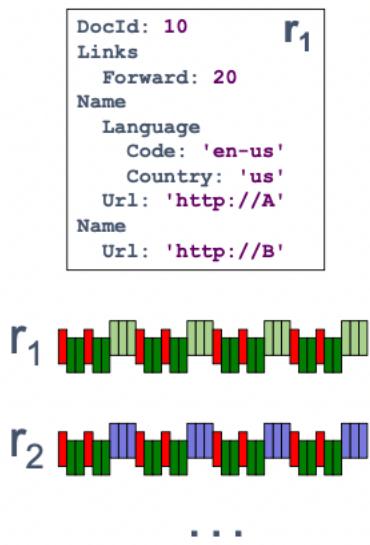
### Big Query System

- Trillion-record, multi-terabyte datasets at interactive speed
  - Scales to thousands of nodes
  - Fault and straggler tolerant execution
- Nested data model
  - Complex datasets; normalization is prohibitive
  - Columnar storage and processing
- Tree architecture (as in web search)
- Interoperates with Google's data mgmt tools
  - In situ data access (e.g., GFS, Bigtable)
  - MapReduce pipelines

## Widely Used inside Google

- Analysis of crawled web documents
- Tracking install data for applications on Android Market
- Crash reporting for Google products
- OCR results from Google Books
- Spam analysis
- Debugging of map tiles on Google Maps
- Tablet migrations in managed Bigtable instances
- Results of tests run on Google's distributed build system
- Disk I/O statistics for hundreds of thousands of disks
- Resource monitoring for jobs run in Google's data centers
- Symbols and dependencies in Google's codebase

## Records      vs.      columns



Challenge: Preserve structure, reconstruct from a subset of fields.

## Nested Data Model

<http://code.google.com/apis/protocolbuffers>

```
multiplicity:
message Document {
    required int64 DocId;           [1,1]
    optional group Links {
        repeated int64 Backward;    [0,*]
        repeated int64 Forward;
    }
    repeated group Name {
        repeated group Language {
            required string Code;
            optional string Country; [0,1]
        }
        optional string Url;
    }
}
```

DocId: 10	r <sub>1</sub>
Links	
Forward: 20	
Forward: 40	
Forward: 60	
Name	
Language	
Code: 'en-us'	
Country: 'us'	
Language	
Code: 'en'	
Url: 'http://A'	
Name	
Url: 'http://B'	
Name	
Language	
Code: 'en-gb'	
Country: 'gb'	

DocId: 20	r <sub>2</sub>
Links	
Backward: 10	
Backward: 30	
Forward: 80	
Name	
Url: 'http://C'	



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

## Column-striped representation

DocId			Name.Url			Links.Forward			Links.Backward		
value	r	d	value	r	d				value	r	d
10	0	0	http://A	0	2				NULL	0	2
20	0	0	http://B	1	2				40	0	2
			NULL	1	1				80	1	2
			http://C	0	2				80	0	2

Name.Language.Code			Name.Language.Country		
value	r	d	value	r	d
en-us	0	2	us	0	3
en	2	2	NULL	2	2
NULL	1	1	NULL	1	1
en-gb	1	2	gb	1	3
NULL	0	1	NULL	0	1



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

## Repetition and definition levels

r=1    r=2    (non-repeating)

### Name.Language.Code

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2
NULL	0	1

record (r=0) has repeated

Language (r=2) has repeated

```

DocId: 10          r1
Links
  Forward: 20
  Forward: 40
  Forward: 60
Name
  Language
    Code: 'en-us'
    Country: 'us'
  Language
    Code: 'en'
    Url: 'http://A'
Name
  Url: 'http://B'
Name
  Language
    Code: 'en-gb'
    Country: 'gb'

```

r: At what repeated field in the field's path  
the value has repeated

```

DocId: 20          r2
Links
  Backward: 10
  Backward: 30
  Forward: 80
Name
  Url: 'http://C'

```

## Query processing

- Optimized for select-project-aggregate
  - Very common class of interactive queries
  - Single scan
  - Within-record and cross-record aggregation
- Approximations: count(distinct), top-k.
- Joins, temp tables, UDFs/TVFs, etc.

## SQL dialect for nested data

```

SELECT DocId AS Id,
       COUNT(Name.Language.Code) WITHIN Name AS Cnt,
       Name.Url + ',' + Name.Language.Code AS Str
  FROM t
 WHERE REGEXP(Name.Url, '^http') AND DocId < 20;

```

### Output table

Id: 10	t <sub>1</sub>
Name	
Cnt: 2	
Language	
Str: 'http://A,en-us'	
Str: 'http://A,en'	
Name	
Cnt: 0	

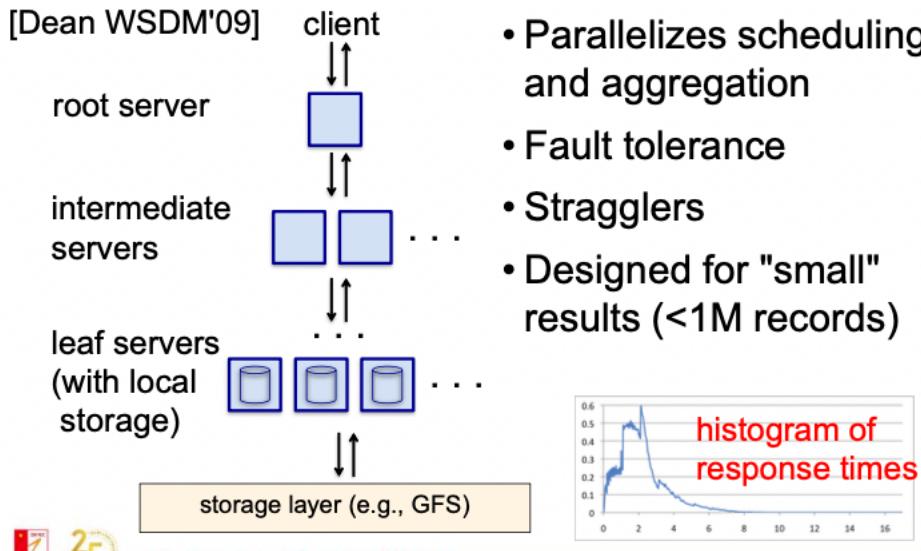
### Output schema

```

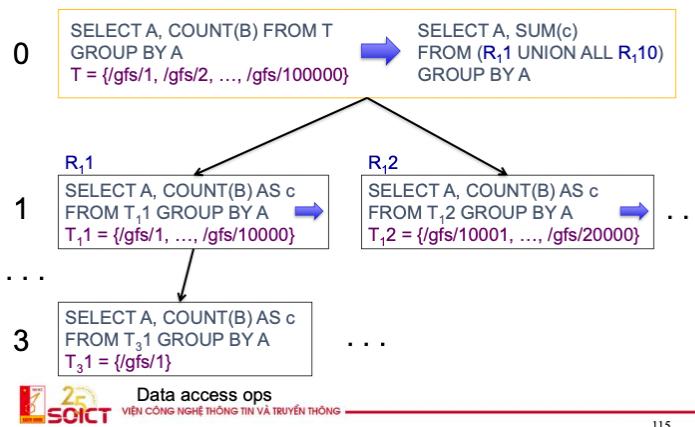
message QueryResult {
  required int64 Id;
  repeated group Name {
    optional uint64 Cnt;
    repeated group Language {
      optional string Str;
    }
  }
}

```

## Serving tree



## Example: count()



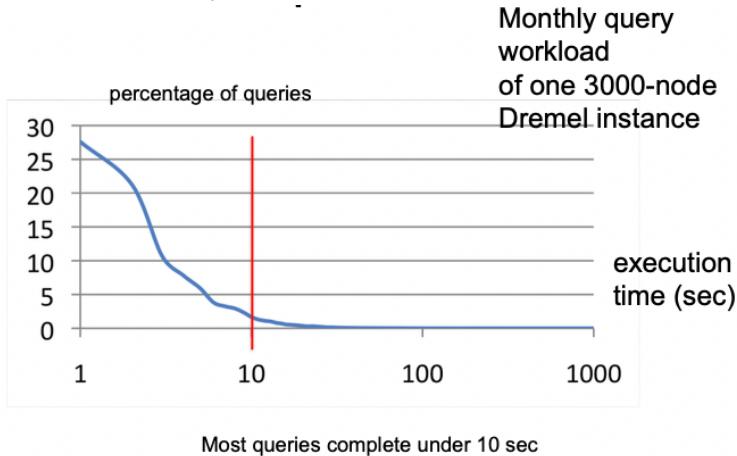
115

## Experiments

- 1 PB of real data (uncompressed, non-replicated).
- 100k-800k tablets per table.
- Experiments run during business hours:

Table name	Number of records	Size (unrepl., compressed)	Number of fields	Data center	Repl. factor
T1	85 billion	87 TB	270	A	3 ×
T2	24 billion	13 TB	530	A	3 ×
T3	4 billion	70 TB	1200	A	3 ×
T4	1+ trillion	105 TB	50	B	3 ×
T5	1+ trillion	20 TB	30	B	2 ×

## Interactive Speed



## Big Query: Powered by Dremel

<http://code.google.com/apis/bigquery/>



## List of Column Data Bases

- Vertica/C-Store
- SybaseIQ
- MonetDB
- LucidDB
- HANA
- Google's Dremel
- Parcell -> Redshit (Another Cloud-DB Service)

## Take-home messages

- OLAP
  - Multi-relational Data model
  - Operators
  - SQL
- Data warehouse (architecture, issues, optimizations)
- Join Processing
- Column Stores (Optimized for OLAP workload)

## Table of Contents

<i>Why we still study OLAP/Data Warehouse in BI?</i> .....	1
<i>Highlights</i> .....	1
<i>Let's get back to the root in 70's: Relational Database</i> .....	1
<i>Basic Structure</i> .....	1
<i>Relation Schema</i> .....	2
<i>Relation Instance</i> .....	2
<i>Database</i> .....	2
<i>Banking Example</i> .....	3
<i>Relational Algebra</i> .....	3
<i>What happens next?</i> .....	3
<i>In early 90's: OLAP &amp; Data Warehouse</i> .....	3
<i>Database Workloads</i> .....	3
<i>OLTP</i> .....	4
<i>OLAP</i> .....	4
<i>One Database or Two?</i> .....	4
<i>OLTP/OLAP Architecture</i> .....	4
<i>OLTP/OLAP Integration</i> .....	4
<i>The Data Warehouse</i> .....	5
<i>Warehouse Architecture</i> .....	5
<i>Star Schemas</i> .....	5
<i>Example: Star Schema</i> .....	5
<i>Visualization – Star Schema</i> .....	6
<i>Dimension and Dependent Attributes</i> .....	6
<i>Warehouse Models &amp; Operators</i> .....	6
<i>Star</i> .....	7
<i>Star Schema</i> .....	7
<i>Terms</i> .....	7
<i>Dimension Hierarchies</i> .....	8
<i>Aggregates</i> .....	8
<i>ROLAP VS. MOLAP</i> .....	9
<i>Cube</i> .....	9
<i>3-D Cube</i> .....	9
<i>Multidimensional Data</i> .....	9
<i>A Sample Data Cube</i> .....	10
<i>Cuboids Corresponding to the Cube</i> .....	10

<i>Cube Aggregation</i>	10
<i>Cube Operators</i>	11
<i>Extended Cube</i>	11
<i>Aggregation Using Hierarchies</i>	11
<i>Pivoting</i>	12
<i>CUBE Operator (SQL-99)</i>	12
<i>Aggregates</i>	12
<i>Query &amp; Analysis Tools</i>	12
<i>Other Operations</i>	13
<i>Data Warehouse Implementation</i>	13
<i>Multi-Tiered Architecture</i>	13
<i>Monitoring</i>	13
<i>Data Cleansing</i>	14
<i>Loading data</i>	14
<i>OLAP Implementation</i>	14
<i>Derived Data</i>	14
<i>What to Materialize</i>	14
<i>Materialization Factors</i>	14
<i>Cube Aggregates Lattice</i>	15
<i>Dimension Hierarchies</i>	15
<i>Interesting Hierarchy</i>	16
<i>Indexing OLAP Data: Bitmap Index</i>	16
<i>Join Processing</i>	16
<i>Schema for Examples</i>	16
<i>Equality Joins with One Join Column</i>	17
<i>Simple Nested Loops Join</i>	17
<i>Block Nested Loops Join</i>	17
<i>Examples of Block Nested Loops</i>	18
<i>Index Nested Loops Join</i>	18
<i>Examples of Index Nested Loops</i>	18
<i>Sort-Merge Join</i>	19
<i>Example of Sort-Merge Join</i>	19
<i>Refinement of Sort-Merger Join</i>	19
<i>Hash-Join</i>	20
<i>Observations in Hash-Join</i>	20
<i>Cost of Hash-Join</i>	20

<i>Join Indices</i> .....	20
<i>General Join Conditions</i> .....	21
<i>An invention in 2000s: Column Stores for OLAP</i> .....	21
<i>Row Store and Column Store</i> .....	21
<i>Column Stores: High Level</i> .....	22
<i>Data Model (Vertica/C-Store)</i> .....	22
<i>Read store: Column Encoding/Compression</i> .....	23
<i>Compression</i> .....	24
<i>Write Store</i> .....	24
<i>How to solve read/write conflict</i> .....	24
<i>Query Execution – Operators</i> .....	24
<i>Benefits in Query Processing</i> .....	24
<i>Evaluation</i> .....	24
<i>Query Performance</i> .....	25
<i>Summary: The performance gain</i> .....	25
<i>Google's Dremel/Big Query: Interactive Analysis of Web-Scale Datasets</i> .....	25
<i>Big Query System</i> .....	25
<i>Widely Used inside Google</i> .....	26
<i>Nested Data Model</i> .....	27
<i>Column-striped representation</i> .....	27
<i>Repetition and definition levels</i> .....	28
<i>Query processing</i> .....	28
<i>SQL dialect for nested data</i> .....	28
<i>Serving tree</i> .....	29
<i>Example: count()</i> .....	29
<i>Experiments</i> .....	29
<i>Interactive Speed</i> .....	30
<i>Big Query: Powered by Dremel</i> .....	30
<i>List of Column Data Bases</i> .....	30
<i>Take-home messages</i> .....	30