



1




2

Agenda

- History of Spark
- Introduction
- Components of Stack
- Resilient Distributed Dataset – RDD

3

History of Spark



4

History of Spark

circa 1979 – **Stanford, MIT, CMU**, etc.
set/list operations in LISP, Prolog, etc., for parallel processing
www-formal.stanford.edu/jmc/history/lisp/lisp.htm

circa 2004 – **Google**
MapReduce: Simplified Data Processing on Large Clusters Jeffrey Dean and Sanjay Ghemawat
research.google.com/archive/mapreduce.html

circa 2006 – **Apache**
Hadoop, originating from the Nutch Project Doug Cutting
research.yahoo.com/files/cutting.pdf

circa 2008 – **Yahoo**
web scale search indexing Hadoop Submit, HUG, etc.
developer.yahoo.com/hadoop/


circa 2009 – **Amazon AWS**
Elastic MapReduce
Hadoop modified for EC2/S3, plus support for Hive, Pig, Cascading, etc.
aws.amazon.com/elasticmapreduce/



5

MapReduce

Most current cluster programming models are based on *acyclic data flow* from stable storage to stable storage



6

MapReduce

- Acyclic data flow is inefficient for applications that repeatedly reuse a *working set* of data:
 - **Iterative** algorithms (machine learning, graphs)
 - **Interactive** data mining tools (R, Excel, Python)



7

Data Processing Improvement Goals



- **Low latency (interactive) queries on historical data:** enable faster decisions

- E.g., identify why a site is slow and fix it



- **Low latency queries on live data (streaming):** enable decisions on real-time data

- E.g., detect & block worms in real-time (a worm may infect **1mil** hosts in **1.3sec**)



- **Sophisticated data processing:** enable “better” decisions

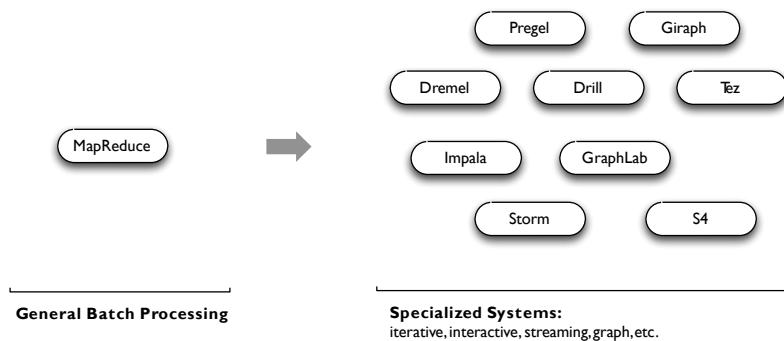
- E.g., anomaly detection, trend analysis

Therefore, people built specialized
systems as workarounds...



8

Specialized Systems



The State of Spark, and Where We're Going Next

Matei Zaharia

Spark Summit (2013)

youtu.be/nU6vO2EJAb4



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

9

Storage vs Processing Wars

NoSQL battles

Relational vs NoSQL

HBase vs Cassandra

Redis vs Memcached vs Riak

MongoDB vs CouchDB vs Couchbase

Neo4j vs Titan vs Giraph vs OrientDB

Solr vs Elasticsearch

Compute battles

MapReduce vs Spark

Spark Streaming vs Storm

Hive vs Spark SQL vs Impala

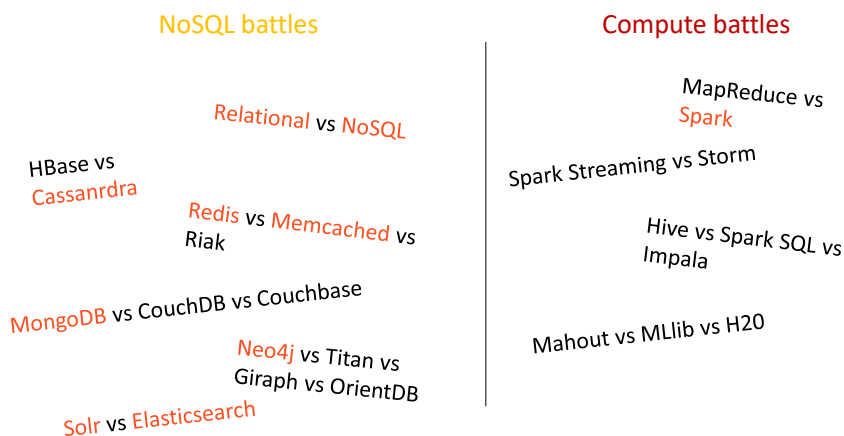
Mahout vs MLlib vs H2O



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

10

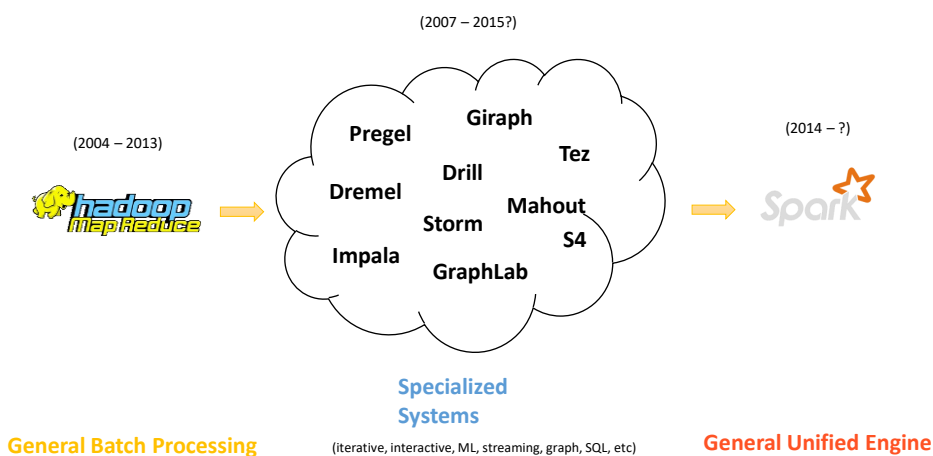
Storage vs Processing Wars



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

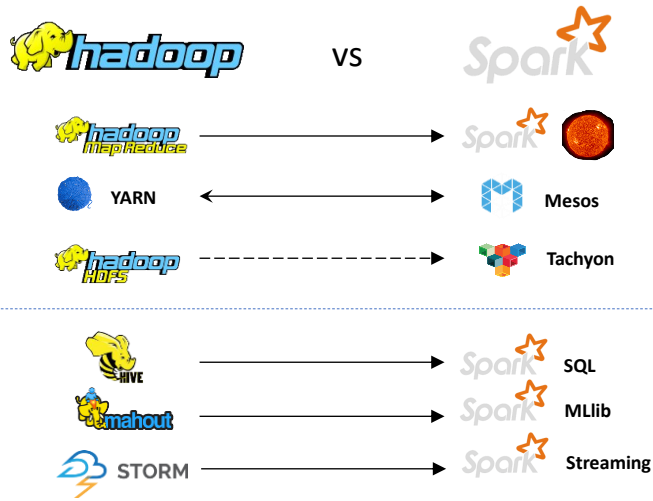
11

Specialized Systems



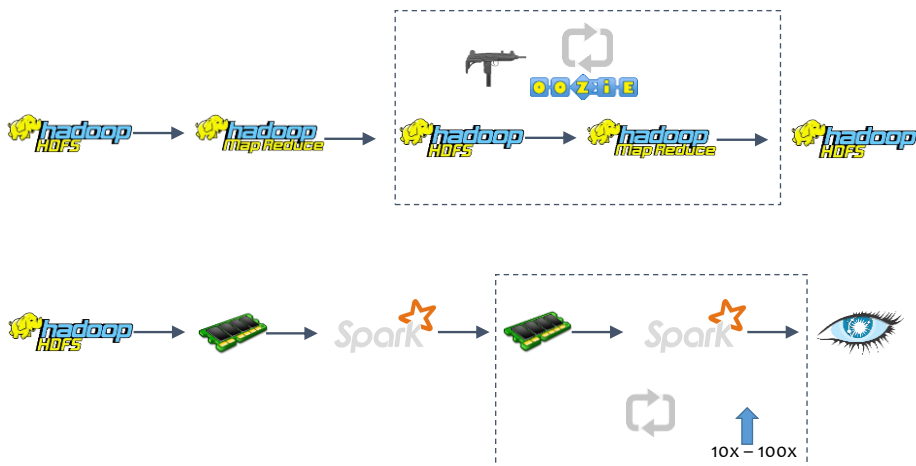
VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

12



13


Spark



14

Support Interactive and Streaming Comp.


- Aggressive use of **memory**
- Why?
 1. Memory transfer rates \gg disk or SSDs
 - Inputs of over 90% of jobs in Facebook, Yahoo!, and Bing clusters fit into memory
 - e.g., 1TB = 1 billion records @ 1KB each
 3. Memory density (still) grows with Moore's law
 - RAM/SSD hybrid memories at horizon



15

Support Interactive and Streaming Comp.


- Increase **parallelism**
- Why?
 - Reduce work per node \rightarrow improve latency
- Techniques:
 - Low latency parallel **scheduler** that achieve high locality
 - Optimized **parallel communication patterns** (e.g., shuffle, broadcast)
 - Efficient **recovery** from failures and straggler mitigation



16

Berkeley AMPLab

- “Launched” January 2011: 6 Year Plan
- 8 CS Faculty
- ~40 students
- 3 software engineers
- Organized for collaboration:



17


Berkeley AMPLab

- Funding:
 - DARPA KData, NSF IRISE Expedition Grant
 - Industrial, founding sponsors
 - 18 other sponsors, including



Goal: Next Generation of Analytics Data Stack for Industry & Research:

- Berkeley Data Analytics Stack (BDAS)
- Release as Open Source

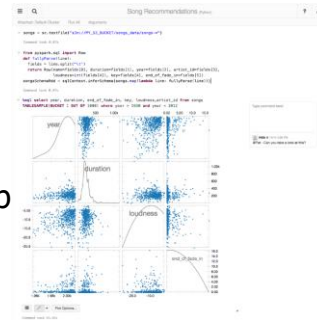


18

Databricks



- Founded in late 2013
- by the creators of Apache Spark
- Original team from UC Berkeley AMPLab
- Raised \$47 Million in 2 rounds



Databricks Cloud:
"A unified platform for building Big Data pipelines – from ETL to Exploration and Dashboards, to Advanced Analytics and Data Products."




19

The Databricks team contributed more than **75%** of the code added to Spark in the 2014



20

History of Spark



Spark: Cluster Computing with Working Sets

Matei Zaharia, Mosharaf Chowdhury,
Michael J. Franklin, Scott Shenker, Ion Stoica
USENIX HotCloud (2010)

people.csail.mit.edu/matei/papers/2010/hotcloud_spark.pdf

Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing

Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave,
Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker,
Ion Stoica NSDI (2012)

usenix.org/system/files/conference/nsdi12/nsdi12-final138.pdf

21

History of Spark



“We present Resilient Distributed Datasets (RDDs), a distributed memory abstraction that lets programmers **perform in-memory computations on large clusters in a fault-tolerant manner**.

RDDs are motivated by two types of applications that current computing frameworks handle inefficiently: **iterative algorithms and interactive data mining tools**.

In both cases, keeping data in memory can **improve performance by an order of magnitude**.”


22

History of Spark

RDD Fault Tolerance

RDDs track the series of transformations used to build them (their *lineage*) to recompute lost data

E.g: `messages = textFile(...).filter(_.contains("error")).map(_.split('\t')(2))`



The State of Spark, and Where We're Going Next

Matei Zaharia

Spark Summit (2013)

[youtu.be/nU6yQ2EIAh4](https://www.youtube.com/watch?v=nU6yQ2EIAh4)



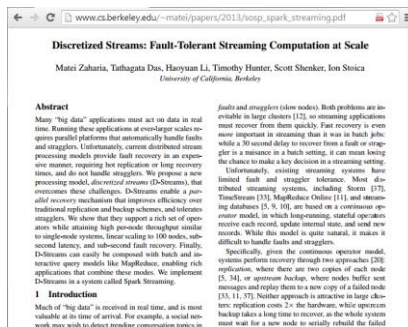
VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

23

History of Spark

Spark **STREAMING**

Analyze real time streams of data in 1/2 second intervals



```

TwitterUtils.createStream(...)
    .filter(_.getText().contains("Spark"))
    .countByWindow(Seconds(5))
  
```



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

24

History of Spark



Seemlessly mix SQL queries with Spark programs.

Spark SQL: Relational Data Processing in Spark

Michael Armbrust¹, Reynold S. Xin¹, Cheng Li¹, Yin Hua¹, Davies Liu¹, Joseph K. Bradley¹,
Xiangrui Meng¹, Tomer Kaitani¹, Michael J. Franklin¹, Ali Ghodsi¹, Matei Zaharia²

¹Databricks Inc. ²MIT CSAIL ³AMPLab, UC Berkeley

ABSTRACT

Spark SQL is a new module in Apache Spark that integrates relational processing with Spark's functional programming API. Built on our experience with Spark, Spark SQL lets Spark programs reuse the benefits of relational processing (e.g., declarative queries and optimized execution), and lets SQL users call complex analytics libraries in Spark (e.g., machine learning). Compared to previous systems, Spark SQL makes two main additions. First, it offers much tighter integration between relational and procedural processing, through a declarative DataFrame API that integrates with procedural Spark code. Second, it includes a highly extensible optimizer Catalyst, built using features of the Scala programming language, that makes it easy to add composable rules, control code generation, and define extension points. Using Catalyst, we have built a variety of features (e.g., schema inference for JDBC, machine learning types, and query federation to external databases) tailored for the complex needs of modern data analysis. We use Spark SQL as an extension of both SQL-on-Spark and of Spark itself, offering richer APIs and optimizations while keeping the benefits of the Spark programming model.

Categories and Subject Descriptors

H.2 [Database Management]: Systems

Keywords

Databases; Data Warehouses; Machine Learning; Spark; Hadoop

1 Introduction

Big data applications require a mix of processing techniques, data sources and storage formats. The real-world systems designed for these workloads, such as MapReduce, were never a unified, but

While the popularity of relational systems shows that users often prefer writing declarative queries, the relational approach is insufficient for many big data applications. First, users want to perform ETLs to and from various data sources that might be semi- or unstructured, requiring custom code. Second, users want to perform advanced analytics, such as machine learning and graph processing, that are challenging to express in relational systems. In practice, we have observed that most data pipelines would ideally be composed with a combination of both relational queries and complex procedural algorithms. Unfortunately, there are two classes of systems—relational and procedural—but still none combined largely disjoint, forcing users to choose one paradigm or the other.

This paper describes our effort to combine both models in Spark SQL, a major new component in Apache Spark 1.0. Spark SQL builds on our earlier SQL-on-Spark effort, called Shark. Rather than forcing users to pick between a relational or a procedural API, however, Spark SQL lets users seamlessly intermix the two.

Spark SQL bridges the gap between the two models through two contributions. First, Spark SQL provides a DataFrame API that can perform relational operations on both external data sources and Spark's built-in distributed collections. This API is similar to the widely used data frame concept in R [13], but extends operations further so that it can perform relational optimizations. Second, to support the wide range of data sources and algorithms in big data, Spark SQL introduces a novel extensible optimizer called Catalyst, which makes it easy to add data sources, optimizer rules, and data types for domains such as machine learning.

The DataFrame API offers rich rich relational/procedural integration within Spark programs. DataFrames are collections of structured records that can be manipulated using Spark's procedural APIs, or using new relational APIs that allow richer optimizations. They can

```
sqlCtx = new HiveContext(sc)
results = sqlCtx.sql(
  "SELECT * FROM people")
names = results.map(lambda p:
  p.name)
```



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

25

History of Spark



Analyze networks of nodes and edges using graph processing

GraphX: A Resilient Distributed Graph System on Spark

Reynold S. Xin, Joseph E. Gonzalez, Michael J. Franklin, Ion Stoica

AMPLab, EECS, UC Berkeley
{xin, jgozalca, franklin, stoica}@cs.berkeley.edu

ABSTRACT

From social networks to targeted advertising, big graphs capture the structure in data and are central to recent advances in machine learning and data mining. Unfortunately, directly applying existing data-parallel tools to graph computation tasks can be cumbersome and inefficient. The need for intuitive, scalable tools for graph computation has led to the development of new graph-parallel systems (e.g., Pregel, PowerGraph) which are designed to efficiently express graph algorithms. Unfortunately, these new graph-parallel systems do not address the challenges of graph construction and transformation which are often just as problematic as the subsequent computation. Furthermore, existing graph-parallel systems provide limited fault-tolerance and support for interactive data mining.

We introduce GraphX, which addresses the challenges of both data-parallel and graph-parallel systems by efficiently expressing graph computation within the Spark data-parallel framework. We leverage new ideas in distributed graph representation to efficiently distribute graphs as similar data structures. Similarly, we leverage advances in data flow systems to explore in-memory computation and fault tolerance. We provide powerful new operators to express graph construction and transformation. Using these primitives we implement the PregelGraph and Pregel algorithms in less than 20 lines of code. Finally, by exploiting the Scala foundation of Spark, we enable users to interactively load, transform, and compute on massive graphs.

1. INTRODUCTION

From social networks to advertising and the web, big graphs can be found in a wide range of important applications. By modeling the

and distributed systems. By abstracting away the challenges of large-scale distributed system design, these frameworks simplify the design, implementation, and application of new sophisticated graph algorithms on large-scale real-world graph problems.

While existing graph-parallel frameworks share many commonalities, they also exhibit a highly different view of graph computation tailored to either the originating domain or a specific family of graph algorithms and applications. Unfortunately, because each framework relies on a separate runtime, it is difficult to compare these abstractions. Furthermore, while these frameworks address the challenges of graph computation, they do not address the challenges of graph construction and transformation in the presence of interacting and applying the results of computation. Finally, few frameworks have built-in support for interactive graph computation.

Alternatively, data-oriented systems like MapReduce and Spark [12] are designed for reliable data processing and are well suited to the task of graph construction (ETL). By exploiting data-parallelism, these systems are highly scalable and support a range of fault-tolerance strategies. However, naively expressing graph computation and graph algorithms in these data-parallel abstractions can be challenging and typically leads to complex joins and recursive data movement that does not exploit the graph structure.

To address these challenges we introduce GraphX, a graph computation system which runs on the Spark data-parallel framework. GraphX extends Spark's Resilient Distributed Dataset (RDD) abstraction to introduce the Resilient Distributed Graph (RDG), which associates vertices with vertices and edges in a graph and provides a collection of expressive computational primitives. Using these

```
graph = Graph(vertices, edges)
messages =
  spark.textFile("hdfs://...")
graph2 =
  graph.joinVertices(messages) {
    (id, vertex, msg) => ...
  }
```


https://amplab.cs.berkeley.edu/wp-content/uploads/2013/05/grades-graphx_with_fonts.pdf



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

26

History of Spark



SQL queries with Bounded Errors and Bounded Response Times

BlinkDB: Queries with Bounded Errors and Bounded Response Times on Very Large Data

Sameer Agarwal¹, Barzan Mozafari², Aurojit Panda¹, Henry Miller¹, Samuel Madden¹, Ion Stoica^{1*}
¹University of California, Berkeley ²Massachusetts Institute of Technology ^{*}Conviva Inc.
 {sameerag, apanda, henrym, istoca}@cs.berkeley.edu, {barzan, madden}@csail.mit.edu

Abstract

In this paper, we present BlinkDB, a massively parallel, approximate query engine for running interactive SQL queries on large volumes of data. BlinkDB allows users to trade-off query accuracy for response time, enabling interactive queries over massive data by running queries on data samples and presenting results annotated with meaningful error bars. To achieve this, BlinkDB uses two key ideas: (1) an adaptive optimization framework that builds and maintains a set of multi-dimensional stratified samples from original data over time, and (2) a dynamic sample selection strategy that selects an appropriately sized sample based on a query's accuracy or response time requirements. We evaluate BlinkDB against the well-known TPC-H benchmarks and a real-world analytic workload derived from Conviva Inc., a company that manages video distribution over the Internet. Our experiments on a real node cluster show that BlinkDB can answer queries on up to 17 TB of data in less than a second (over 100x faster than Hive), within an error of 2-10%.

1. Introduction

Modern data analytics applications involve computing aggregates over a large number of records to mid-up web clicks,

counting of large amounts of data by trading result accuracy for response time and space. These techniques include sampling [10, 14], sketches [12], and on-line aggregation [6]. To illustrate the utility of such techniques, consider the following simple query that computes the average SessionTime over all users originating in New York:

```
SELECT AVG(SessionTime)
FROM Sessions
WHERE City = 'New York'
```

Suppose the Sessions table contains 100 million tuples for New York, and cannot fit in memory. In that case, the above query may take a long time to execute, since disk reads are expensive, and such a query would need multiple disk accesses to stream through all the tuples. Suppose we instead executed the same query on a sample containing only 10,000 New York tuples, such that the entire sample fits in memory. This would be orders of magnitude faster, while still providing an approximate result within a few percent of the actual value, an accuracy good enough for many practical purposes. Using sampling theory we could even provide confidence bounds on the accuracy of the answer [16].

Previously described approximation techniques make different trade-offs between efficiency and the generality of the

```
SELECT avg(sessionTime)
FROM Table
WHERE city='San Francisco'
WITHIN 2 SECONDS
```

Queries with Time Bounds

```
SELECT avg(sessionTime)
FROM Table
WHERE city='San Francisco'
ERROR 0.1 CONFIDENCE 95.0%
```

Queries with Error Bounds

https://www.cs.berkeley.edu/~sameerag/blinkdb_eurosys13.pdf



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

27

History of Spark

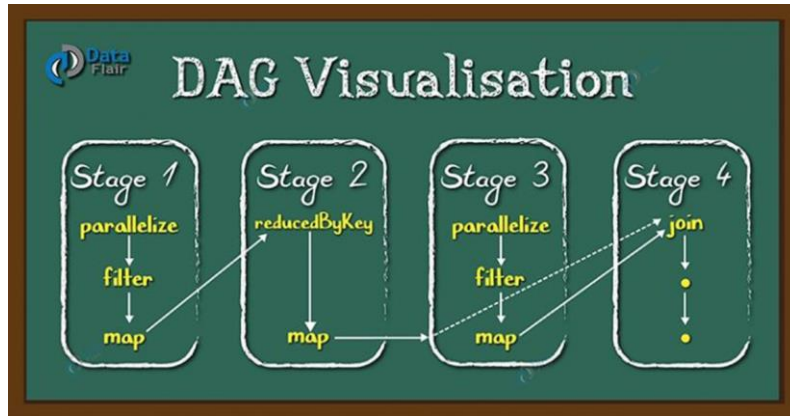
- Unlike the various specialized systems, Spark's goal was to *generalize* MapReduce to support new apps within same engine
- Two reasonably small additions are enough to express the previous models:
 - *fast data sharing*
 - *general DAGs*
- This allows for an approach which is more efficient for the engine, and much simpler for the end users



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG


28

Directed Acyclic Graph - DAG



What is Apache Spark

- Spark is a unified **analytics** engine **for large-scale data processing**
- **Speed:** run workloads 100x faster
 - High performance for both batch and streaming data
 - Computations run in memory



What is Apache Spark

- **Ease of Use:** write applications quickly in Java, Scala, Python, R, SQL
 - Offer over 80 high-level operators
 - Use them interactively from Scala, Python, R, and SQL

```
df = spark.read.json("logs.json")
df.where("age > 21")
select("name.first").show()
```


Spark's Python DataFrame API
Read JSON files with automatic schema inference



31

What is Apache Spark

- **Generality:** combine SQL, Streaming, and complex analytics
 - Provide libraries including SQL and DataFrames, Spark Streaming, MLib, GraphX
 - Wide range of workloads e.g., batch applications, interactive algorithms, interactive queries, streaming



32

What is Apache Spark

• Run Everywhere:



- run on Hadoop, Apache Mesos, Kubernetes, standalone or in the cloud.
- access data in HDFS, Aluxio, Apache Cassandra, Apache Hbase, Apache Hive, etc.



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

33

Comparison between Hadoop and Spark

		
Strengths	<ul style="list-style-type: none"> ▪ Can collect any data ▪ Limitless in size 	<ul style="list-style-type: none"> ▪ Can work off any Hadoop collection ▪ Runs on Hadoop, or other clusters ▪ In-memory processing makes it very fast ▪ Supports Java, Scala, Python, and R*, and can be used with SQL.
Used for	<ul style="list-style-type: none"> ▪ Initial data ingestion ▪ Data curation ▪ Large-scale “boil the ocean” analytics ▪ Data archiving 	<ul style="list-style-type: none"> ▪ Complex query processing of large amounts of data quickly ▪ Can handle ad hoc queries
Limitations	<ul style="list-style-type: none"> ▪ MapReduce is hard to program ▪ Disk-based batch nature limits speed, agility. 	<ul style="list-style-type: none"> ▪ Limited only by processor speed, available memory, cores, and cluster size.



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

34

100TB Daytona Sort Competition

	Hadoop MR Record	Spark Record	Spark 1 PB
Data Size	102.5 TB	100 TB	1000 TB
Elapsed Time	72 mins	23 mins	234 mins
# Nodes	2100	206	190
# Cores	50400 physical	6592 virtualized	6080 virtualized
Cluster disk throughput	3150 GB/s (est.)	618 GB/s	570 GB/s
Sort Benchmark Daytona Rules	Yes	Yes	No
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network	virtualized (EC2) 10Gbps network
Sort rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Sort rate/node	0.67 GB/min	20.7 GB/min	22.5 GB/min

Spark sorted the same data **3X faster** using **10X fewer machines** than Hadoop MR in 2013.

All the sorting took place on disk (HDFS)
without using Spark's in-memory cache!



35

The screenshot shows a Wired article from October 10, 2014, at 2:36 PM. The article is titled "Startup Crunches 100 Terabytes of Data in a Record 23 Minutes" by Clint Finley. It features social media sharing buttons for Facebook, Twitter, and LinkedIn. Below the article title, there is a large yellow graphic with the number "23" and the word "MINUTES". To the right of the graphic, there is a section titled "MUST READS" with three links: "Google launches Contributor, a...", "Net neutrality looks doomed in...", and "Five tech products that...". Below this, there is a section titled "Databricks demolishes big data benchmark to prove Spark is fast on disk, too" by Derrick Harris, dated Oct 10, 2014, at 1:49 PM PST. The article is part of a series by GigaOM, with links to Cloud, Data, Media, Mobile, Science & Energy, Social & Web, and Podcasts. The GigaOM logo is in the top right corner, along with links to SIGN IN and SUBSCRIBE.

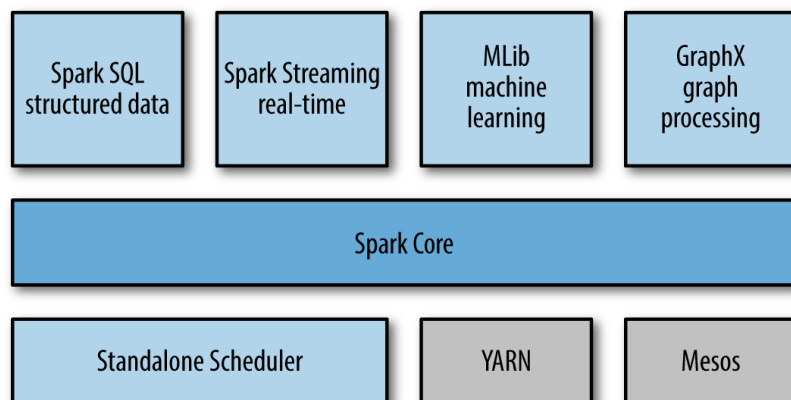


36

Components of Stack

37

The Spark stack



38

The Spark stack

- Spark Core:
 - contain basic functionality of Spark including task scheduling, memory management, fault recovery, etc.
 - provide APIs for building and manipulating RDDs
- SparkSQL
 - allow querying structured data via SQL, Hive Query Language
 - allow combining SQL queries and data manipulations in Python, Java, Scala



39

The Spark stack

- Spark Streaming: enables processing of live streams of data via APIs
- Mlib:
 - contain common machine language functionality
 - provide multiple types of algorithms: classification, regression, clustering, etc.
- GraphX:
 - library for manipulating graphs and performing graph-parallel computations
 - extend Spark RDD API



40

The Spark stack

- Cluster Managers
 - Hadoop Yarn
 - Apache Mesos, and
 - Standalone Scheduler (simple manager in Spark).

Resilient Distributed Dataset – RDD

- **RDD Basics**
- Creating RDDs
- RDD Operations
- Common Transformation and Actions
- Persistence (Caching)

RDD Basics

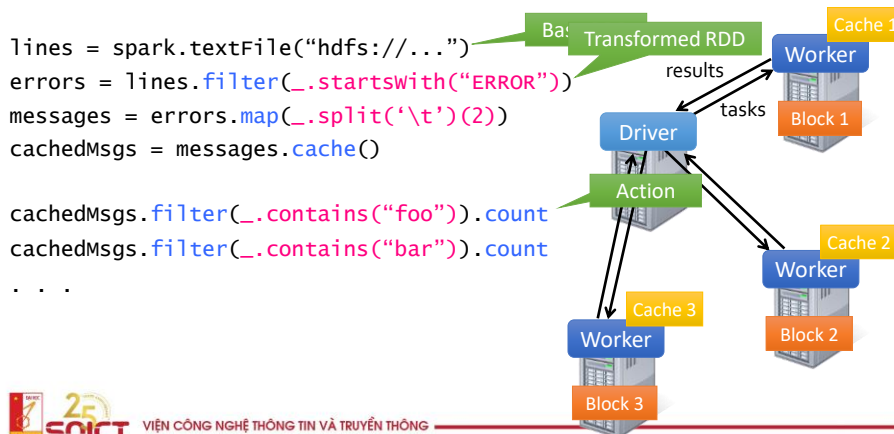
- RDD:
 - Immutable distributed collection of objects
 - Split into multiple partitions => can be computed on different nodes
- All work in Spark is expressed as
 - creating new RDDs
 - transforming existing RDDs
 - calling actions on RDDs



43

Example


Load error messages from a log into memory, then interactively search for various patterns



44

RDD Basics

- Two types of operations: transformations and actions
- Transformations: construct a new RDD from a previous one e.g., filter data
- Actions: compute a result base on an RDD e.g., count elements, get first element



45

Transformations

- Create new RDDs from existing RDDs
- Lazy evaluation
 - See the whole chain of transformations
 - Compute just the data needed
- Persist contents:
 - persist an RDD in memory to reuse it in future
 - persist RDDs on disk is possible

46

Typical works of a Spark program

1. Create some input RDDs form external data
2. Transform them to define new RDDs using transformations like filter()
3. Ask Spark to persist() any intermediate RDDs that will need to be reused
4. Launch actions such as count(), first() to kick off a parallel computation



47

Resilient Distributed Dataset – RDD

- RDD Basics
- **Creating RDDs**
- RDD Operations
- Common Transformation and Actions
- Persistence (Caching)



48

Two ways to create RDDs

1. Parallelizing a collection: uses parallelize()

- Python

```
lines = sc.parallelize(["pandas", "i like pandas"])
```

- Scala

```
val lines = sc.parallelize(List("pandas", "i like pandas"))
```

- Java

```
JavaRDD<String> lines =  
sc.parallelize(Arrays.asList("pandas", "i like  
pandas"));
```



49

Two ways to create RDDs

2. Loading data from external storage

- Python

```
lines = sc.textFile("/path/to/README.md")
```

- Scala

```
val lines = sc.textFile("/path/to/README.md")
```

- Java

```
JavaRDD<String> lines =  
sc.textFile("/path/to/README.md");
```



50

Resilient Distributed Dataset – RDD

- RDD Basics
- Creating RDDs
- **RDD Operations**
- Common Transformation and Actions
- Persistence (Caching)

RDD Operations

- Two types of operations
 - Transformations: operations that **return a new RDDs** e.g., map(), filter()
 - Actions: operations that return a **result** to the driver program or write it to storage such as count(), first()
- Treated differently by Spark
 - Transformation: lazy evaluation
 - Action: execution at any time

Transformation

- Example 1. Use **filter()**

- Python

```
inputRDD = sc.textFile("log.txt")
errorsRDD = inputRDD.filter(lambda x: "error" in x)
```

- Scala

```
val inputRDD = sc.textFile("log.txt")
val errorsRDD = inputRDD.filter(line =>
line.contains("error"))
```

- Java

```
JavaRDD<String> inputRDD = sc.textFile("log.txt");
JavaRDD<String> errorsRDD = inputRDD.filter(
new Function<String, Boolean>() {
public Boolean call(String x) {
return x.contains("error"); }}
});
```



53

Transformation

- **filter()**

- does not change the existing *inputRDD*
- returns a pointer to an entirely new RDD
- *inputRDD* still can be reused

- **union()**

```
errorsRDD = inputRDD.filter(lambda x: "error" in x)
warningsRDD=inputRDD.filter(lambda x: "warning" in x)
badLinesRDD = errorsRDD.union(warningsRDD)
```


- transformations can operate on any number of input RDDs



54

Transformation

- Spark keeps track dependencies between RDDs, called the **lineage graph**
- Allow recovering lost data



Actions

- Example. count the number of errors
- Python

```

print "Input had " + badLinesRDD.count() + " concerning lines"
print "Here are 10 examples:"
for line in badLinesRDD.take(10):
    print line
  
```

- Scala

```

println("Input had " + badLinesRDD.count() + " concerning lines")
println("Here are 10 examples:")
badLinesRDD.take(10).foreach(println)
  
```

- Java

```

System.out.println("Input had " + badLinesRDD.count() + " concerning
lines")
System.out.println("Here are 10 examples:")
for (String line: badLinesRDD.take(10)) {
    System.out.println(line);
}
  
```

Resilient Distributed Dataset – RDD

- RDD Basics
- Creating RDDs
- RDD Operations
- **Common Transformation and Actions**
- Persistence (Caching)



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

57

RDD Basics

Transformations	Actions
map flatMap filter sample union groupByKey reduceByKey join cache	reduce collect count save lookupKey ...



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

58

Transformations

transformation	description
map (<i>func</i>)	return a new distributed dataset formed by passing each element of the source through a function <i>func</i>
filter (<i>func</i>)	return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true
flatMap (<i>func</i>)	similar to map, but each input item can be mapped to 0 or more output items (so <i>func</i> should return a Seq rather than a single item)
sample (<i>withReplacement</i> , <i>fraction</i> , <i>seed</i>)	sample a fraction <i>fraction</i> of the data, with or without replacement, using a given random number generator <i>seed</i>
union (<i>otherDataset</i>)	return a new dataset that contains the union of the elements in the source dataset and the argument
distinct (<i>[numTasks]</i>)	return a new dataset that contains the distinct elements of the source dataset



59

Transformations

transformation	description
groupByKey (<i>[numTasks]</i>)	when called on a dataset of (<i>K</i> , <i>V</i>) pairs, returns a dataset of (<i>K</i> , Seq[<i>V</i>]) pairs
reduceByKey (<i>func</i> , <i>[numTasks]</i>)	when called on a dataset of (<i>K</i> , <i>V</i>) pairs, returns a dataset of (<i>K</i> , <i>V</i>) pairs where the values for each key are aggregated using the given reduce function
sortByKey (<i>[ascending]</i> , <i>[numTasks]</i>)	when called on a dataset of (<i>K</i> , <i>V</i>) pairs where <i>K</i> implements Ordered, returns a dataset of (<i>K</i> , <i>V</i>) pairs sorted by keys in ascending or descending order, as specified in the boolean <i>ascending</i> argument
join (<i>otherDataset</i> , <i>[numTasks]</i>)	when called on datasets of type (<i>K</i> , <i>V</i>) and (<i>K</i> , <i>W</i>), returns a dataset of (<i>K</i> , (<i>V</i> , <i>W</i>)) pairs with all pairs of elements for each key
cogroup (<i>otherDataset</i> , <i>[numTasks]</i>)	when called on datasets of type (<i>K</i> , <i>V</i>) and (<i>K</i> , <i>W</i>), returns a dataset of (<i>K</i> , Seq[<i>V</i>], Seq[<i>W</i>]) tuples – also called <i>groupWith</i>
cartesian (<i>otherDataset</i>)	when called on datasets of types <i>T</i> and <i>U</i> , returns a dataset of (<i>T</i> , <i>U</i>) pairs (all pairs of elements)



60

Actions

<i>action</i>	<i>description</i>
reduce(<i>func</i>)	aggregate the elements of the dataset using a function <i>func</i> (which takes two arguments and returns one), and should also be commutative and associative so that it can be computed correctly in parallel
collect()	return all the elements of the dataset as an array at the driver program – usually useful after a filter or other operation that returns a sufficiently small subset of the data
count()	return the number of elements in the dataset
first()	return the first element of the dataset – similar to <i>take(1)</i>
take(<i>n</i>)	return an array with the first <i>n</i> elements of the dataset – currently not executed in parallel, instead the driver program computes all the elements
takeSample(<i>withReplacement</i>, <i>fraction</i>, <i>seed</i>)	return an array with a random sample of <i>num</i> elements of the dataset, with or without replacement, using the given random number generator seed



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

61

Actions

<i>action</i>	<i>description</i>
saveAsTextFile(<i>path</i>)	write the elements of the dataset as a text file (or set of text files) in a given directory in the local filesystem, HDFS or any other Hadoop-supported file system. Spark will call <i>toString</i> on each element to convert it to a line of text in the file
saveAsSequenceFile(<i>path</i>)	write the elements of the dataset as a Hadoop <i>SequenceFile</i> in a given path in the local filesystem, HDFS or any other Hadoop-supported file system. Only available on RDDs of key-value pairs that either implement Hadoop's <i>Writable</i> interface or are implicitly convertible to <i>Writable</i> (Spark includes conversions for basic types like <i>Int</i> , <i>Double</i> , <i>String</i> , etc).
countByKey()	only available on RDDs of type (K, V) . Returns a <i>Map</i> of (K, Int) pairs with the count of each key
foreach(<i>func</i>)	run a function <i>func</i> on each element of the dataset – usually done for side effects such as updating an accumulator variable or interacting with external storage systems

62

Resilient Distributed Dataset – RDD

- RDD Basics
- Creating RDDs
- RDD Operations
- Common Transformation and Actions
- Persistence (Caching)



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

63

Persistence levels

Level	Space used	CPU time	In memory	On disk	Comments
MEMORY_ONLY	High	Low	Y	N	
MEMORY_ONLY_SER	Low	High	Y	N	
MEMORY_AND_DISK	High	Medium	Some	Some	Spills to disk if there is too much data to fit in memory.
MEMORY_AND_DISK_SER	Low	High	Some	Some	Spills to disk if there is too much data to fit in memory. Stores serialized representation in memory.
DISK_ONLY	Low	High	N	Y	



VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG

64

Persistence

• Example

```
val result = input.map(x => x * x)
result.persist(StorageLevel.DISK_ONLY)
println(result.count())
println(result.collect().mkString(", "))
```



65

Acknowledgement and
References

Books:

- Holden Karau, Andy Konwinski, Patrick Wendell & Matei Zaharia. Learning Spark. Oreilly
- Tutorialspoint. Spark Core Programming

Slides:

- Paco Nathan. Intro to Apache Spark
- Harold Liu. Berkely Data Analytics Stack
- DataBricks. Intro to Spark Development



66

Q&A



67



68



69



70